

Transformers and Pretraining

Computational Linguistics


Alexander Koller

08 December 2023


Recurrent neural networks

- When we process language, we usually don't have inputs of fixed length (sentences can be arbitrarily long).


Tagging



Language modeling

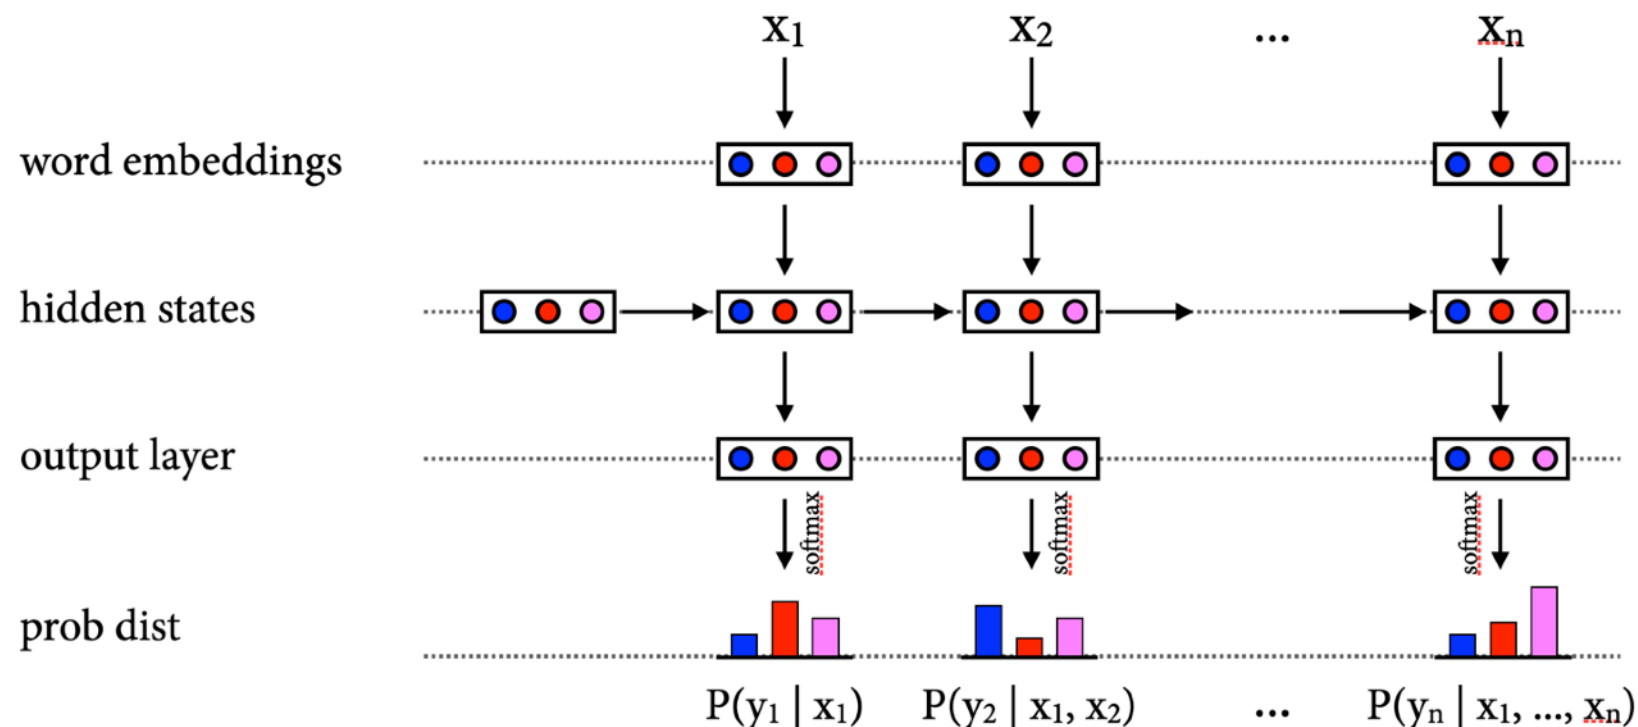


Translation




Limitations of RNNs

- Long computation paths make backprop hard.
- Computing loss at timestep t requires doing all the computations at times $1, \dots, t-1$.
 - ▶ Computations for different timesteps are different.
 - ▶ Can't parallelize these computations.

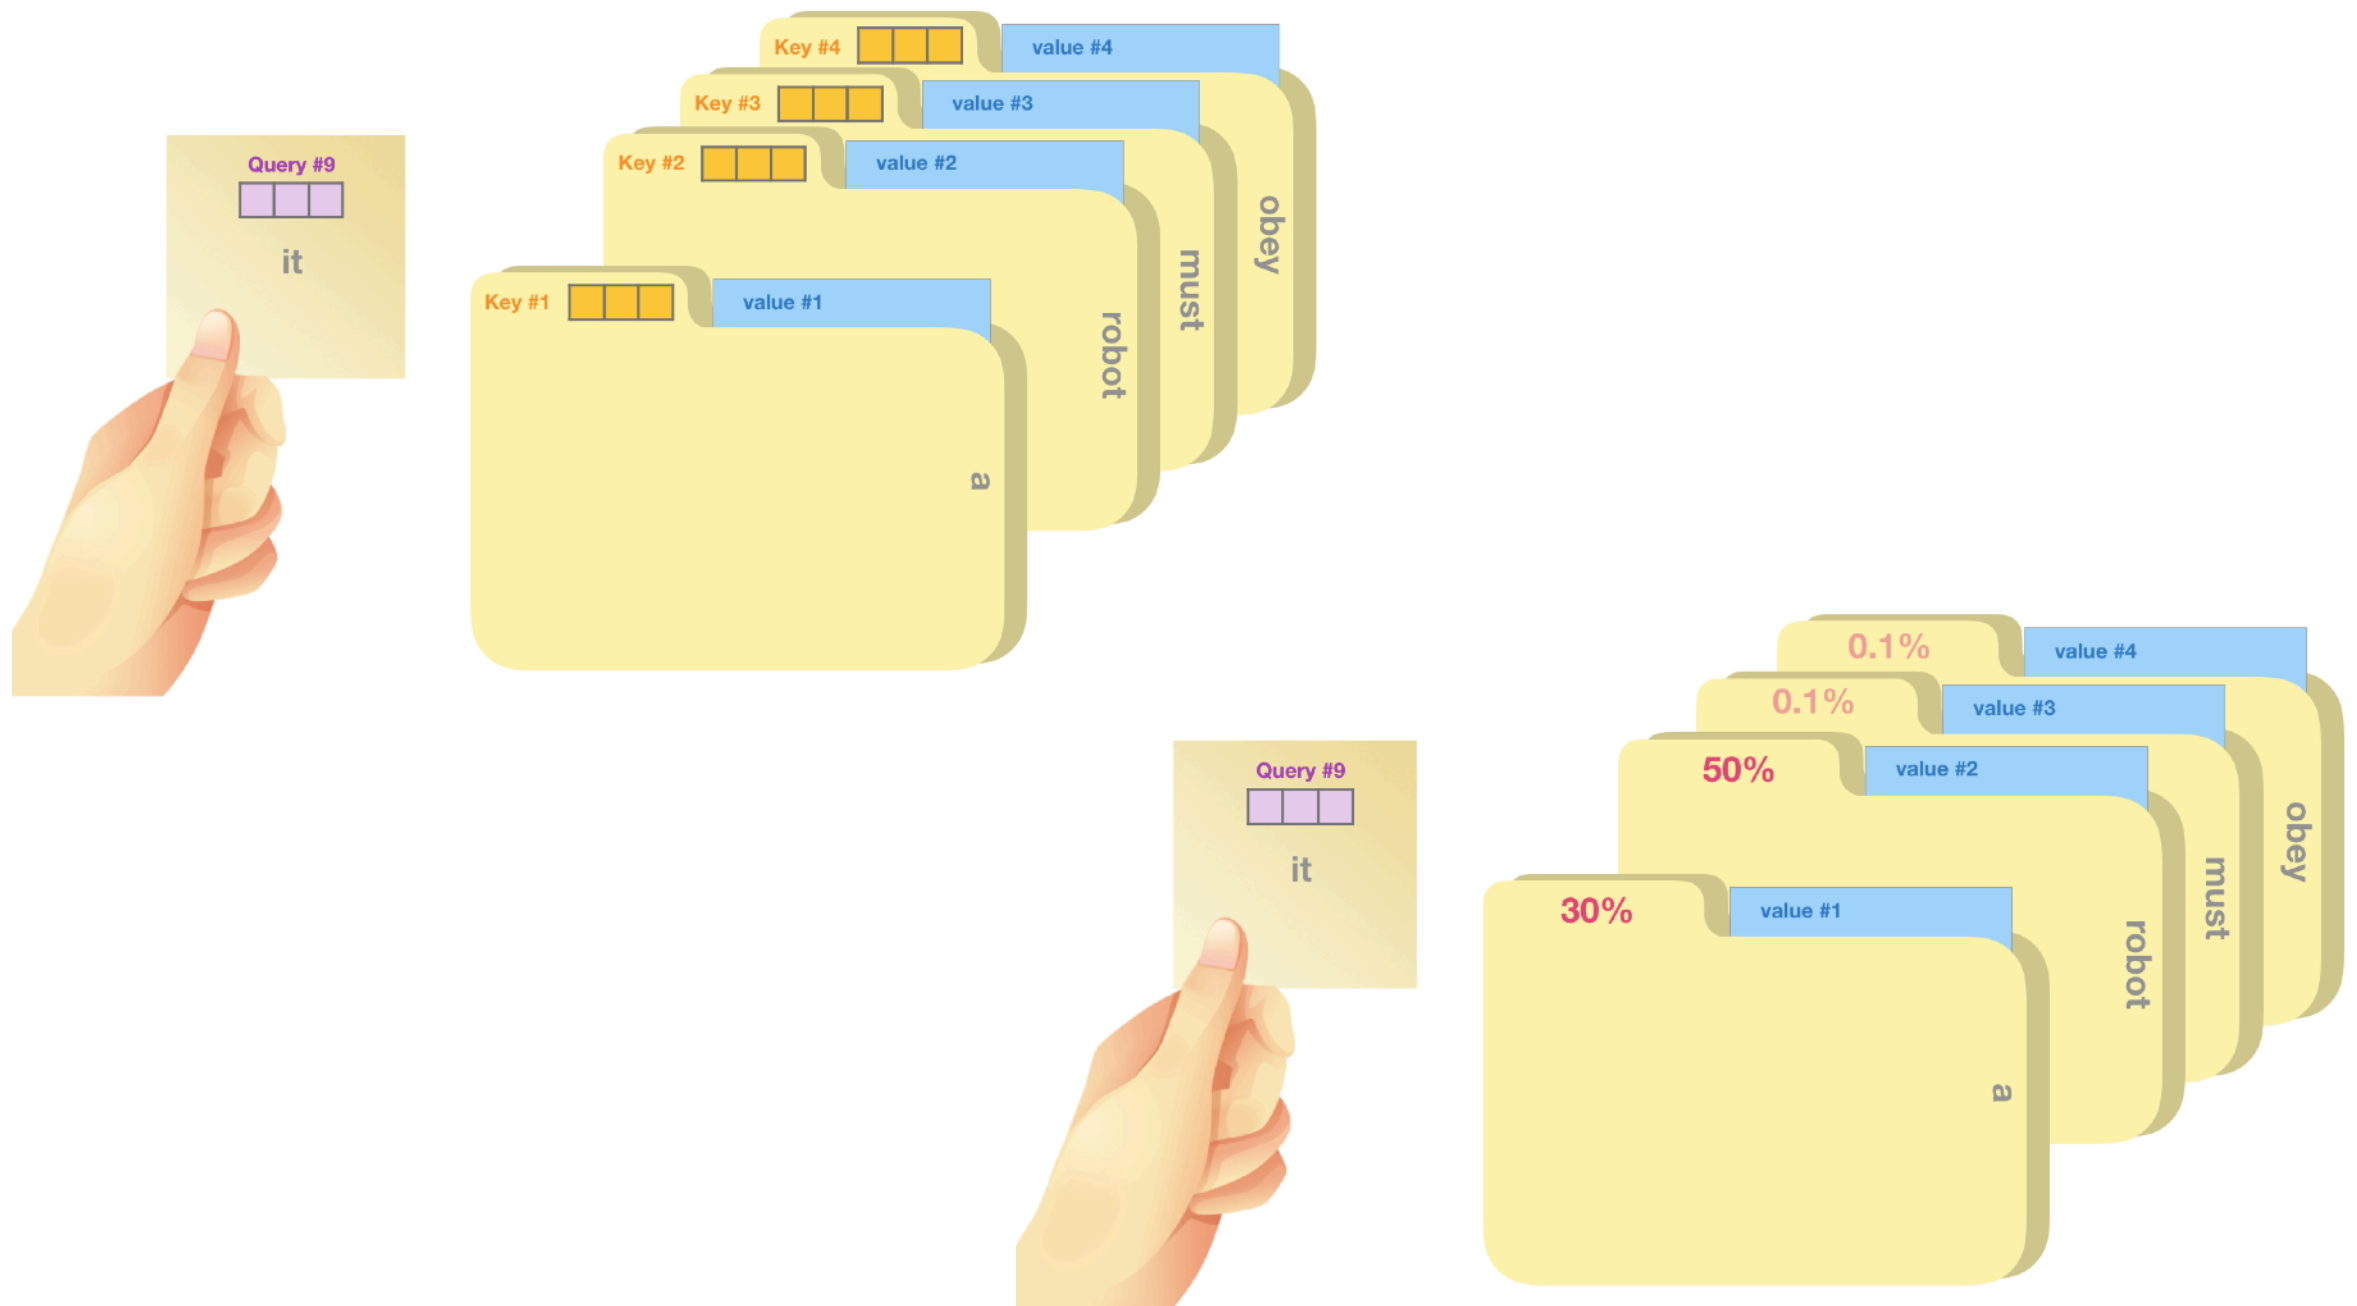


Transformers: Big picture

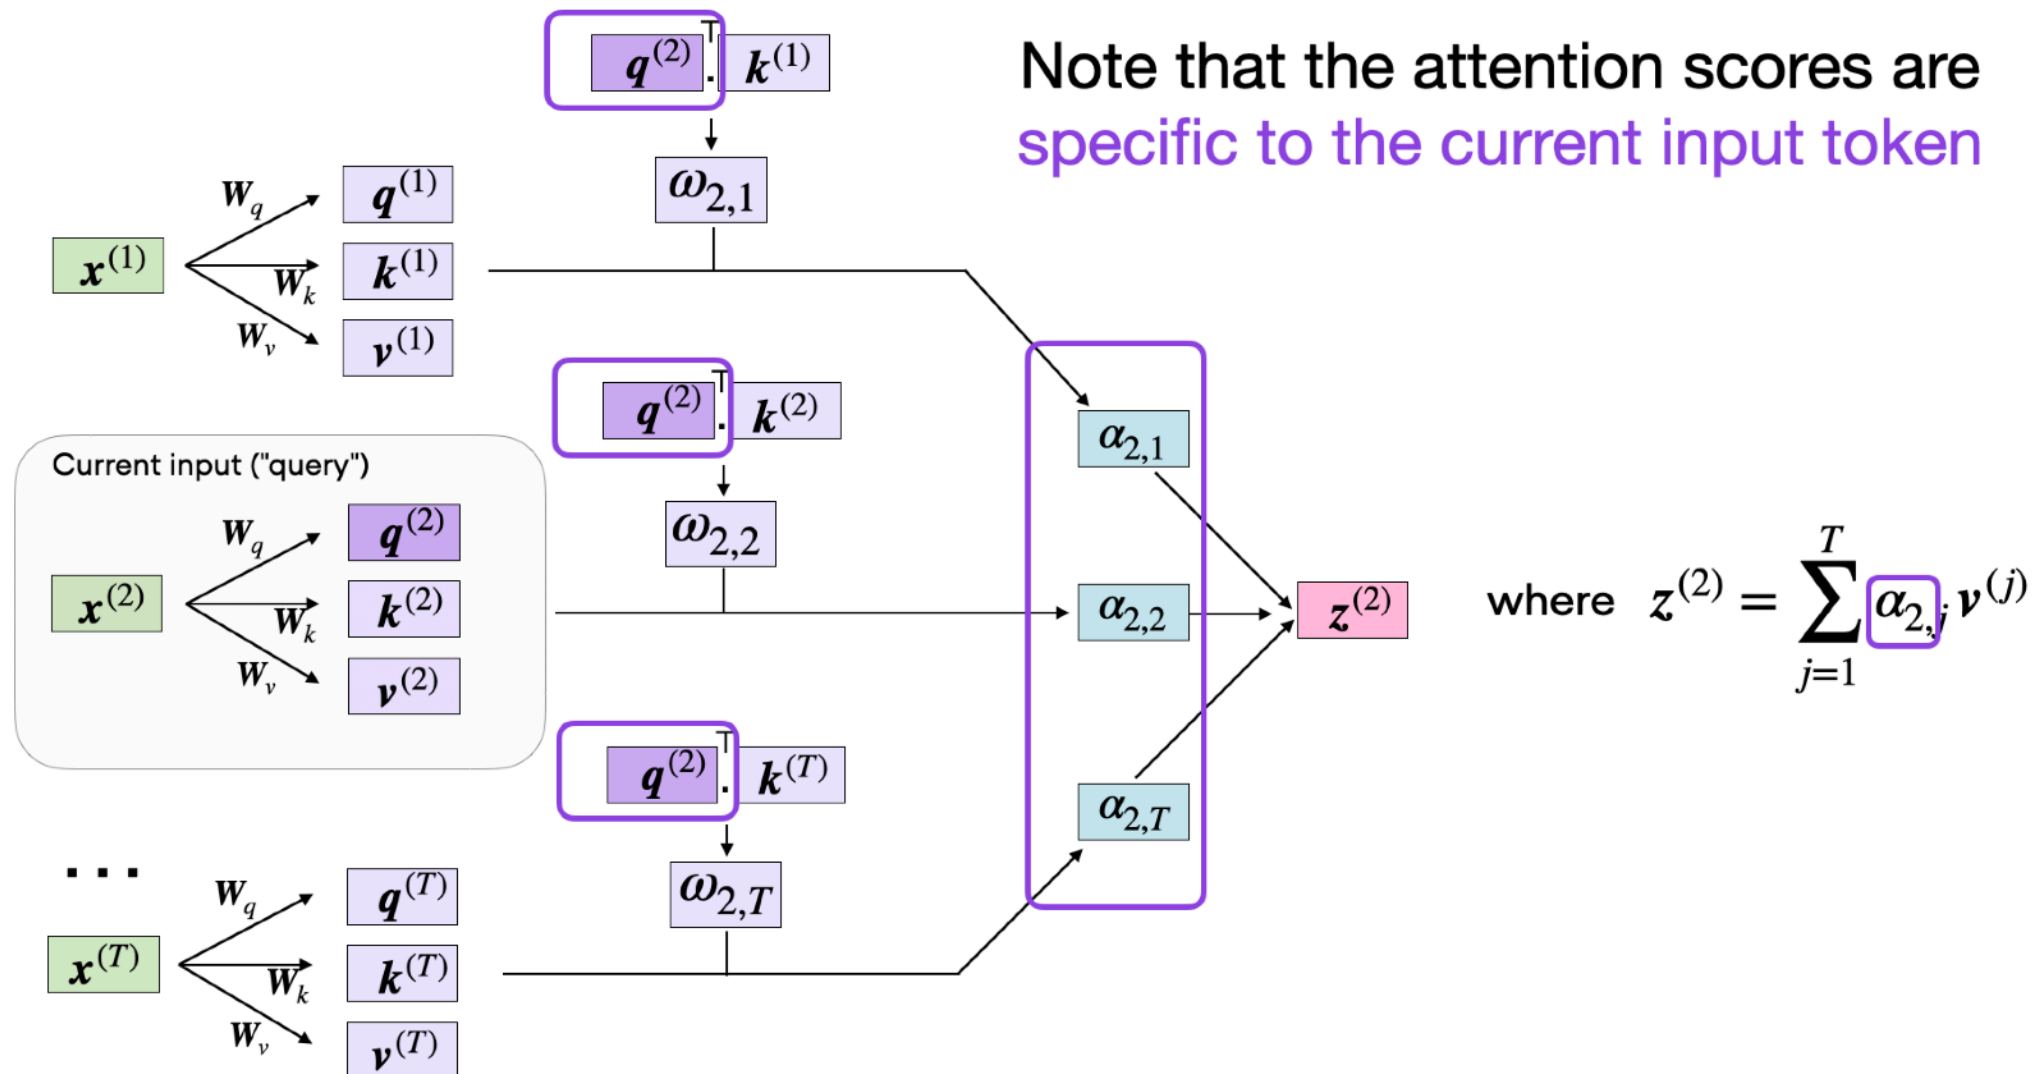
Encode sequence of inputs in parallel!



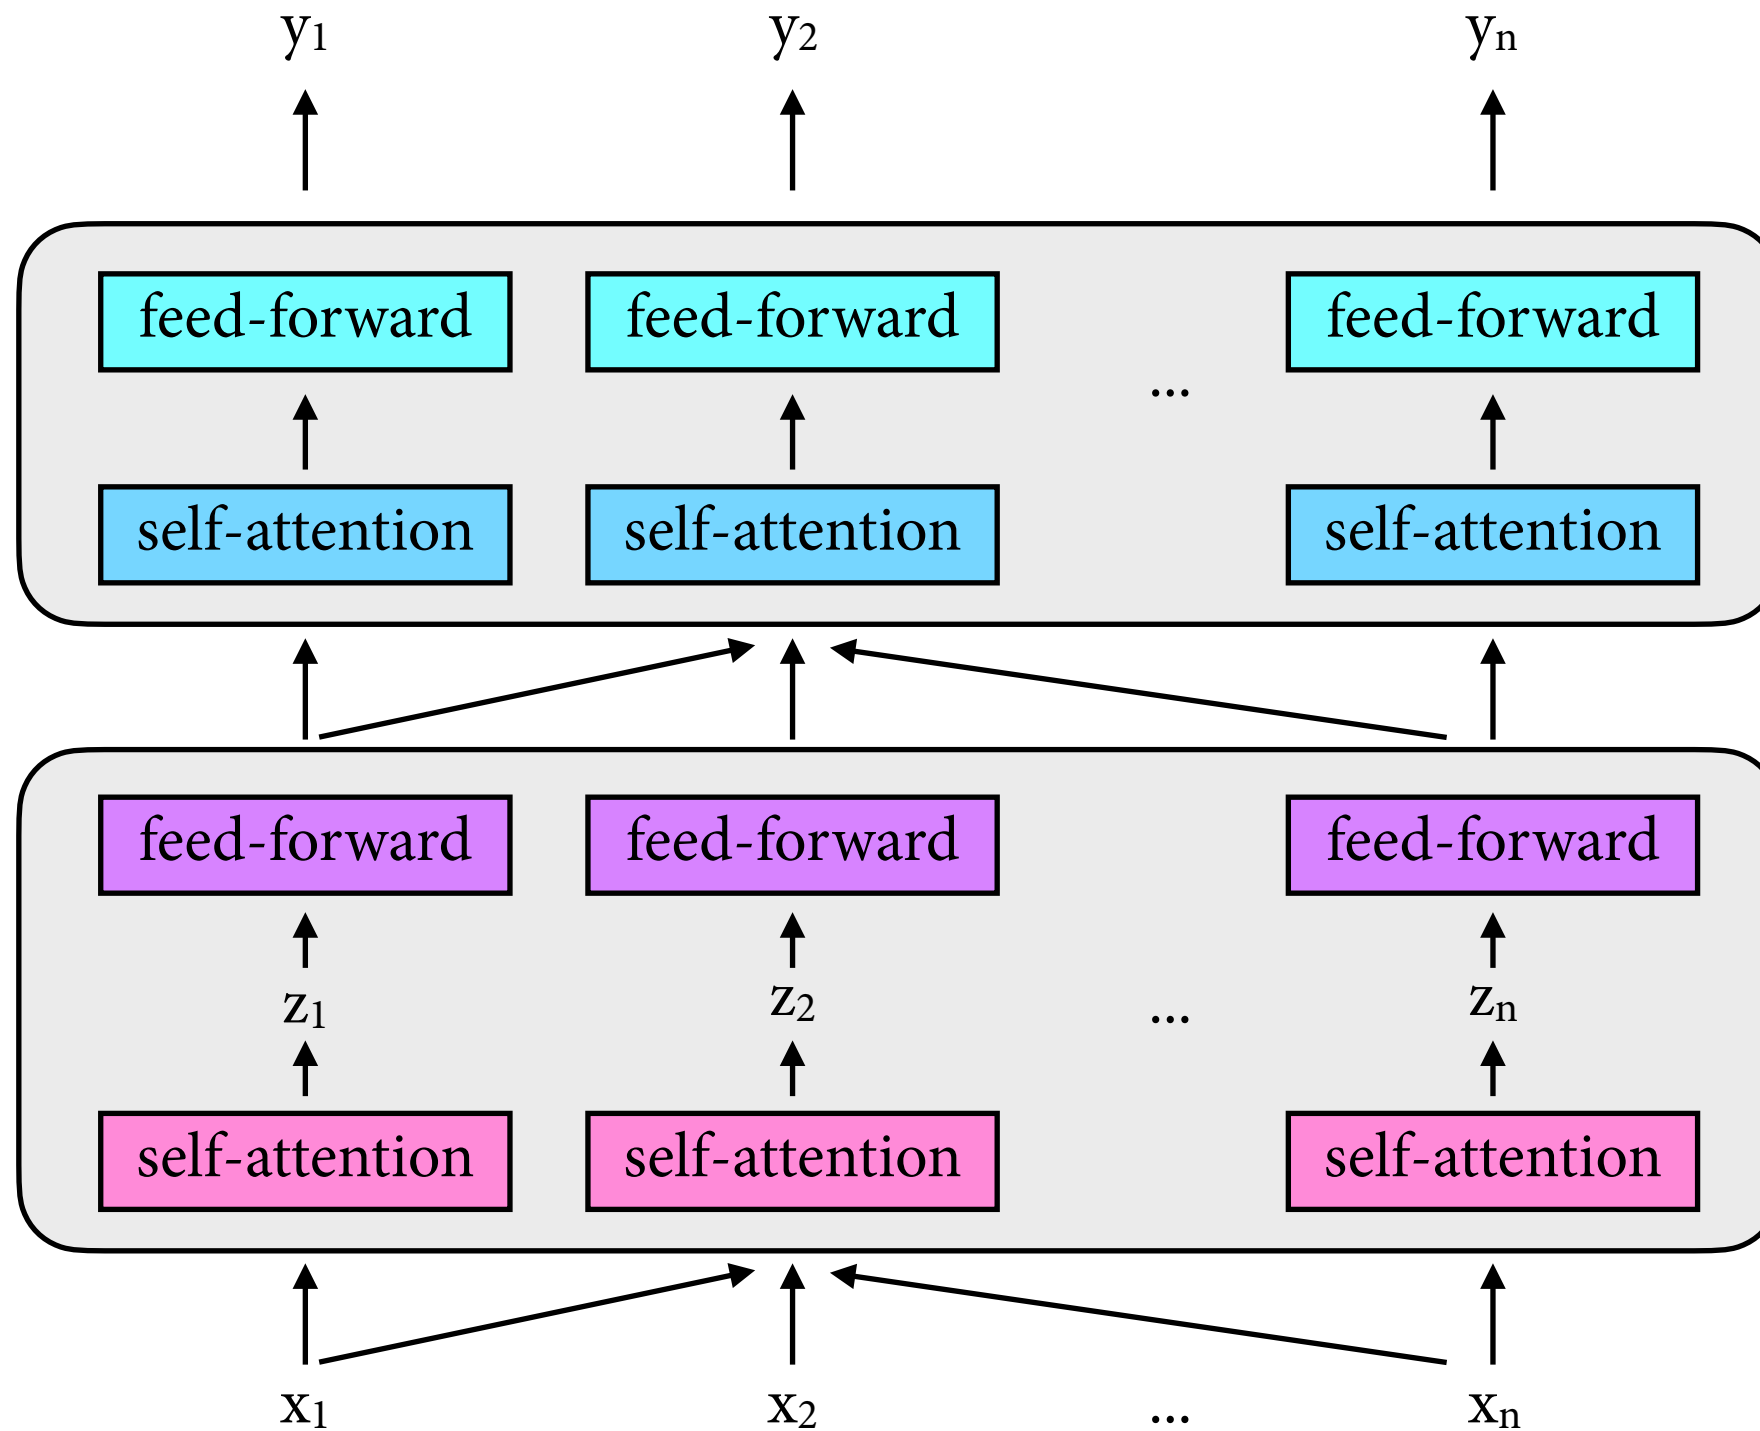
Attention: Big picture



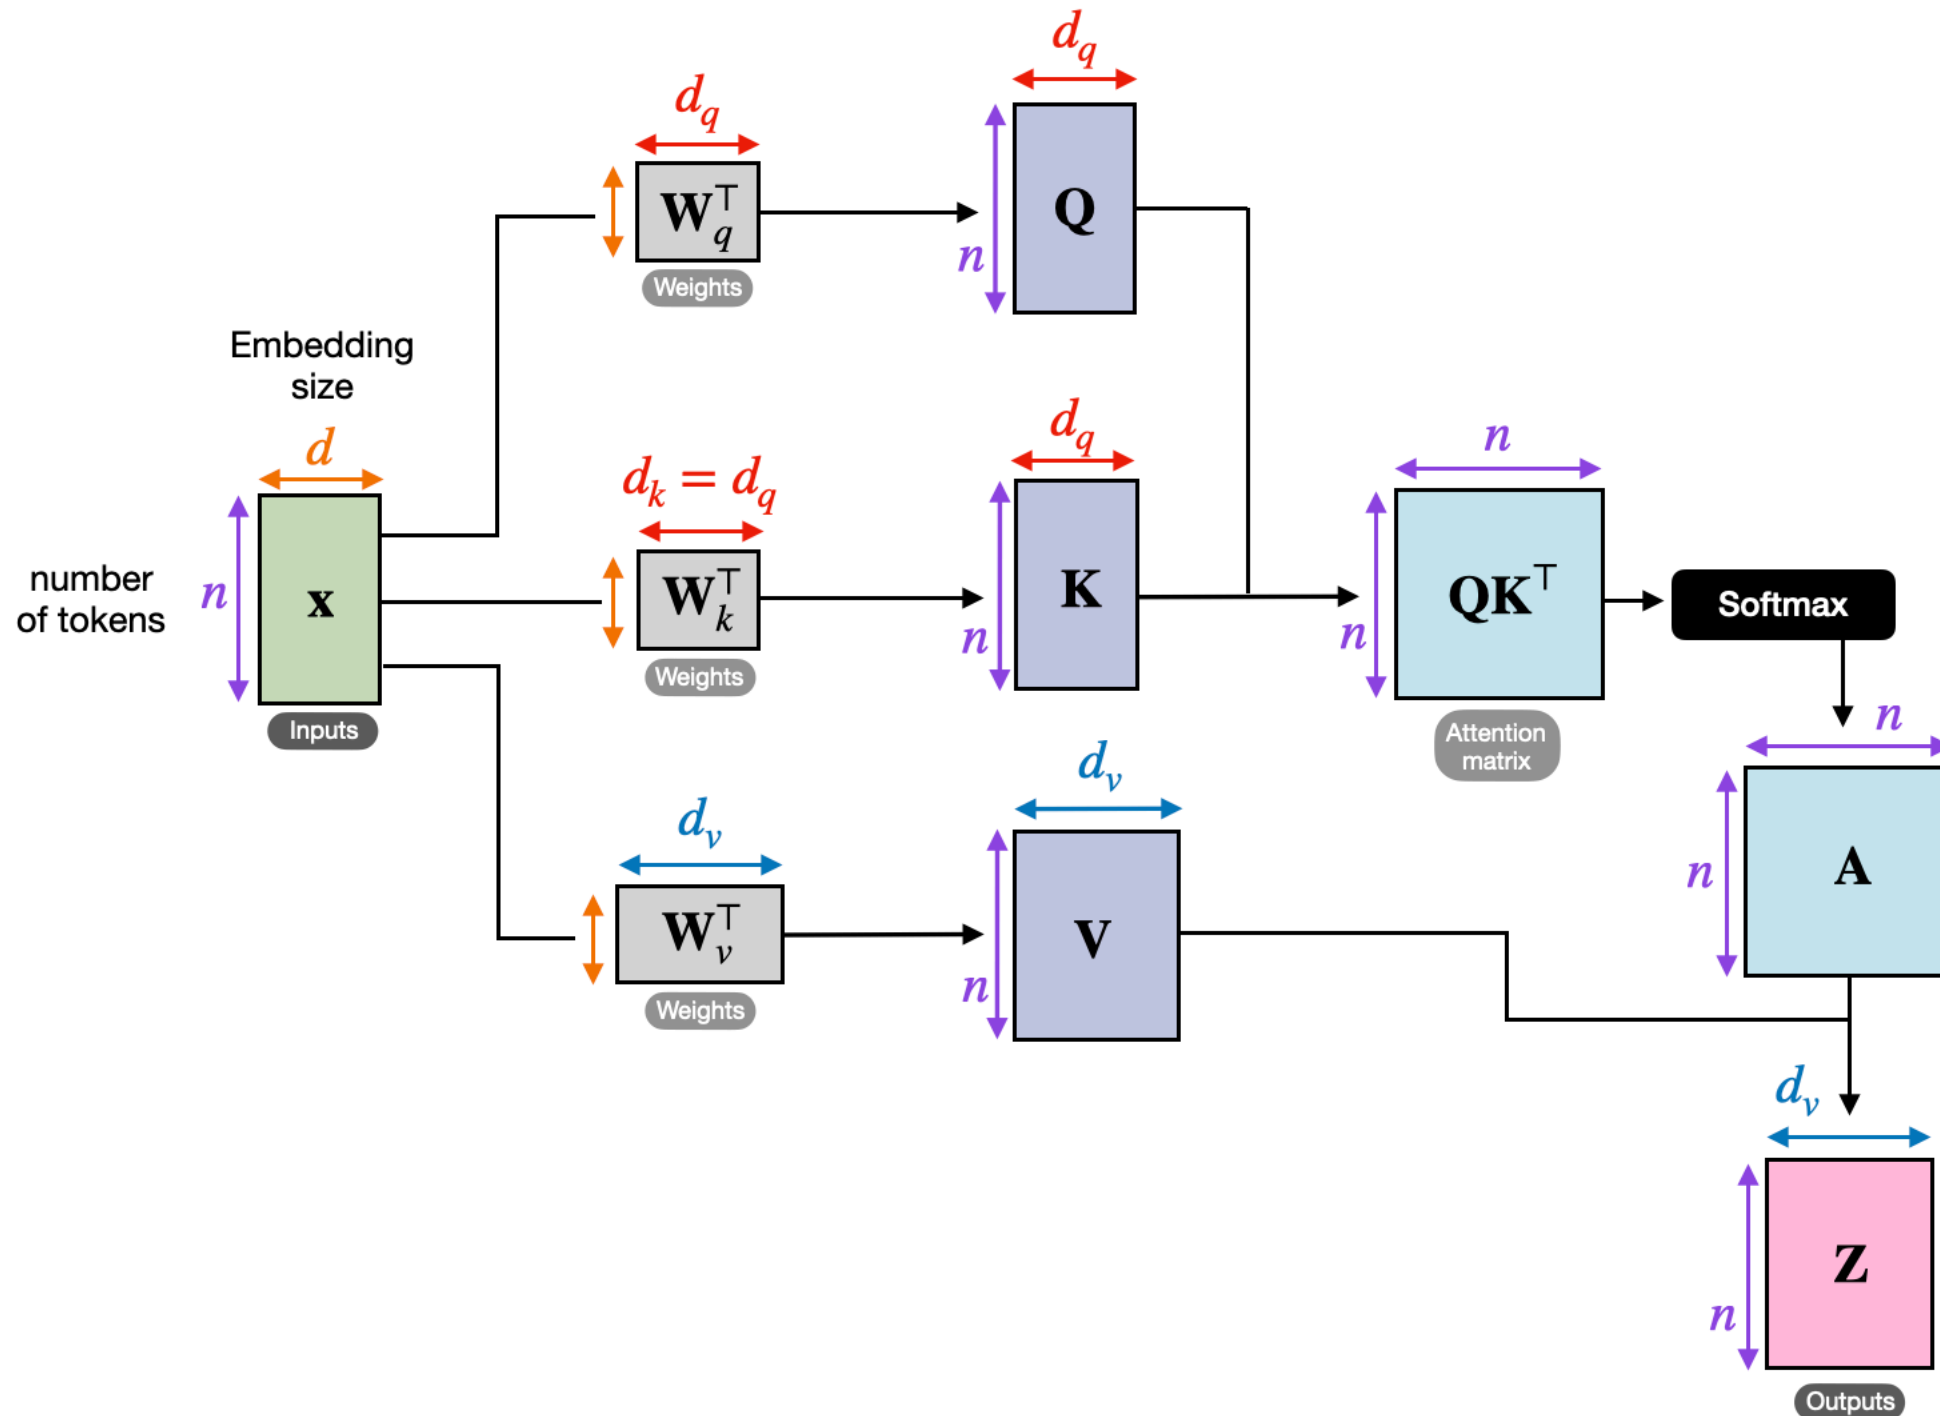
Attention: Details



Transformer Encoder



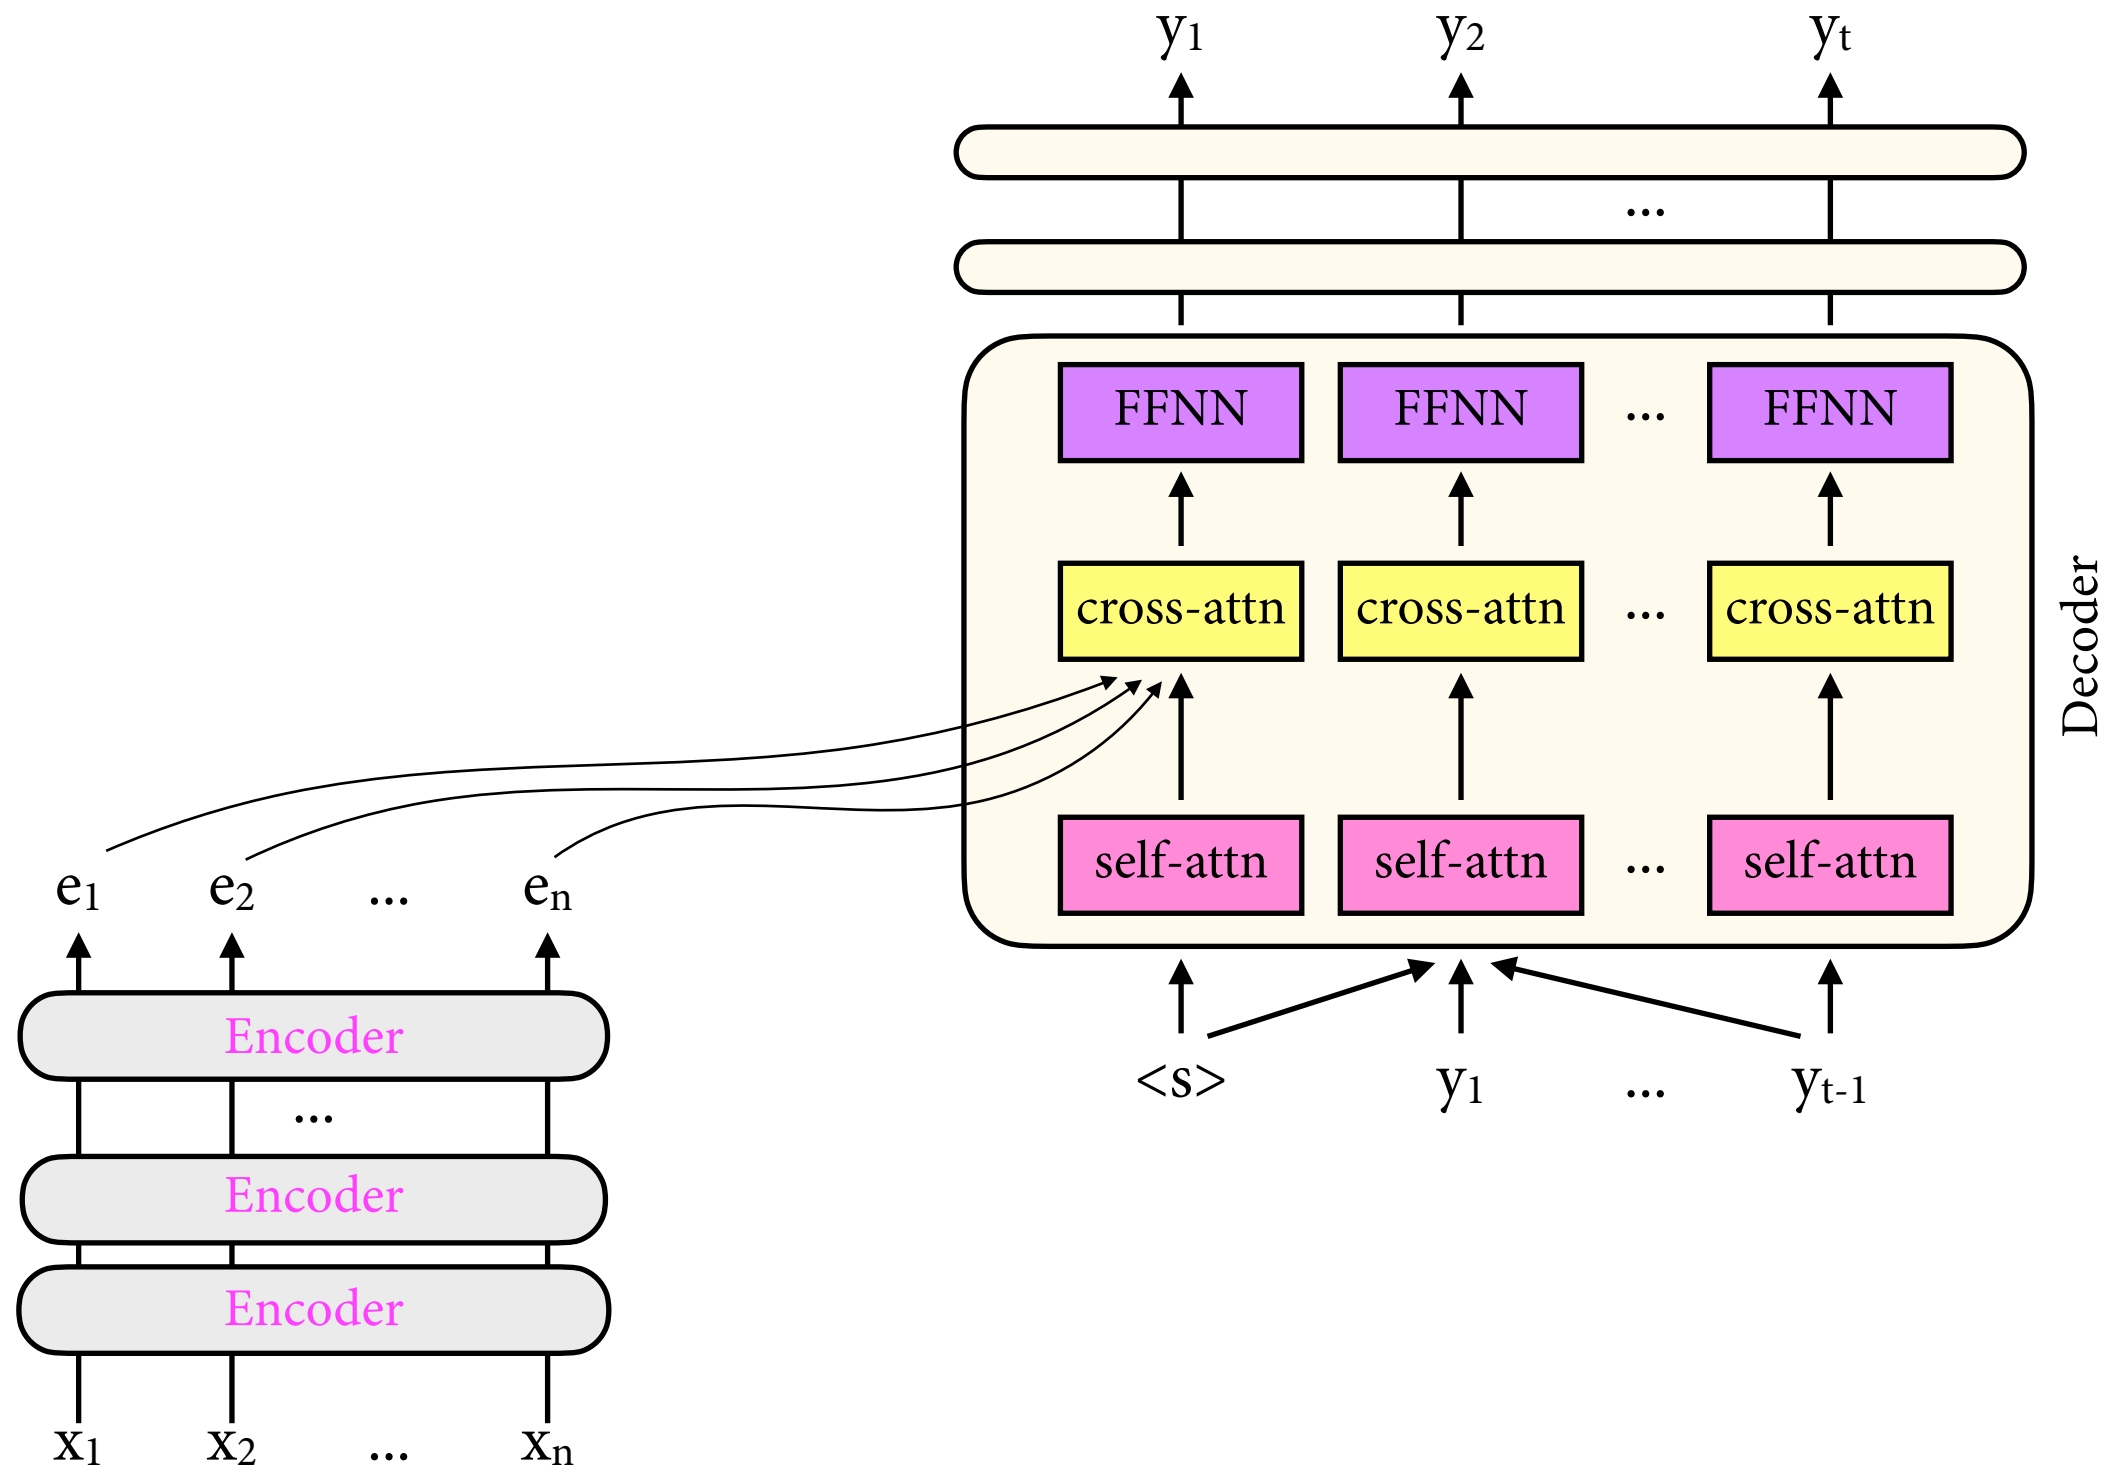
Self-attention dimensions



Some extra details

- Need to add *positional encodings* to each input, so the transformer can keep string positions apart.
- Typical transformers have *multiple attention heads* rather than just one. Can pay independent attention, outputs are concatenated.
- I skipped over some technical details that are important in practice; see original paper.

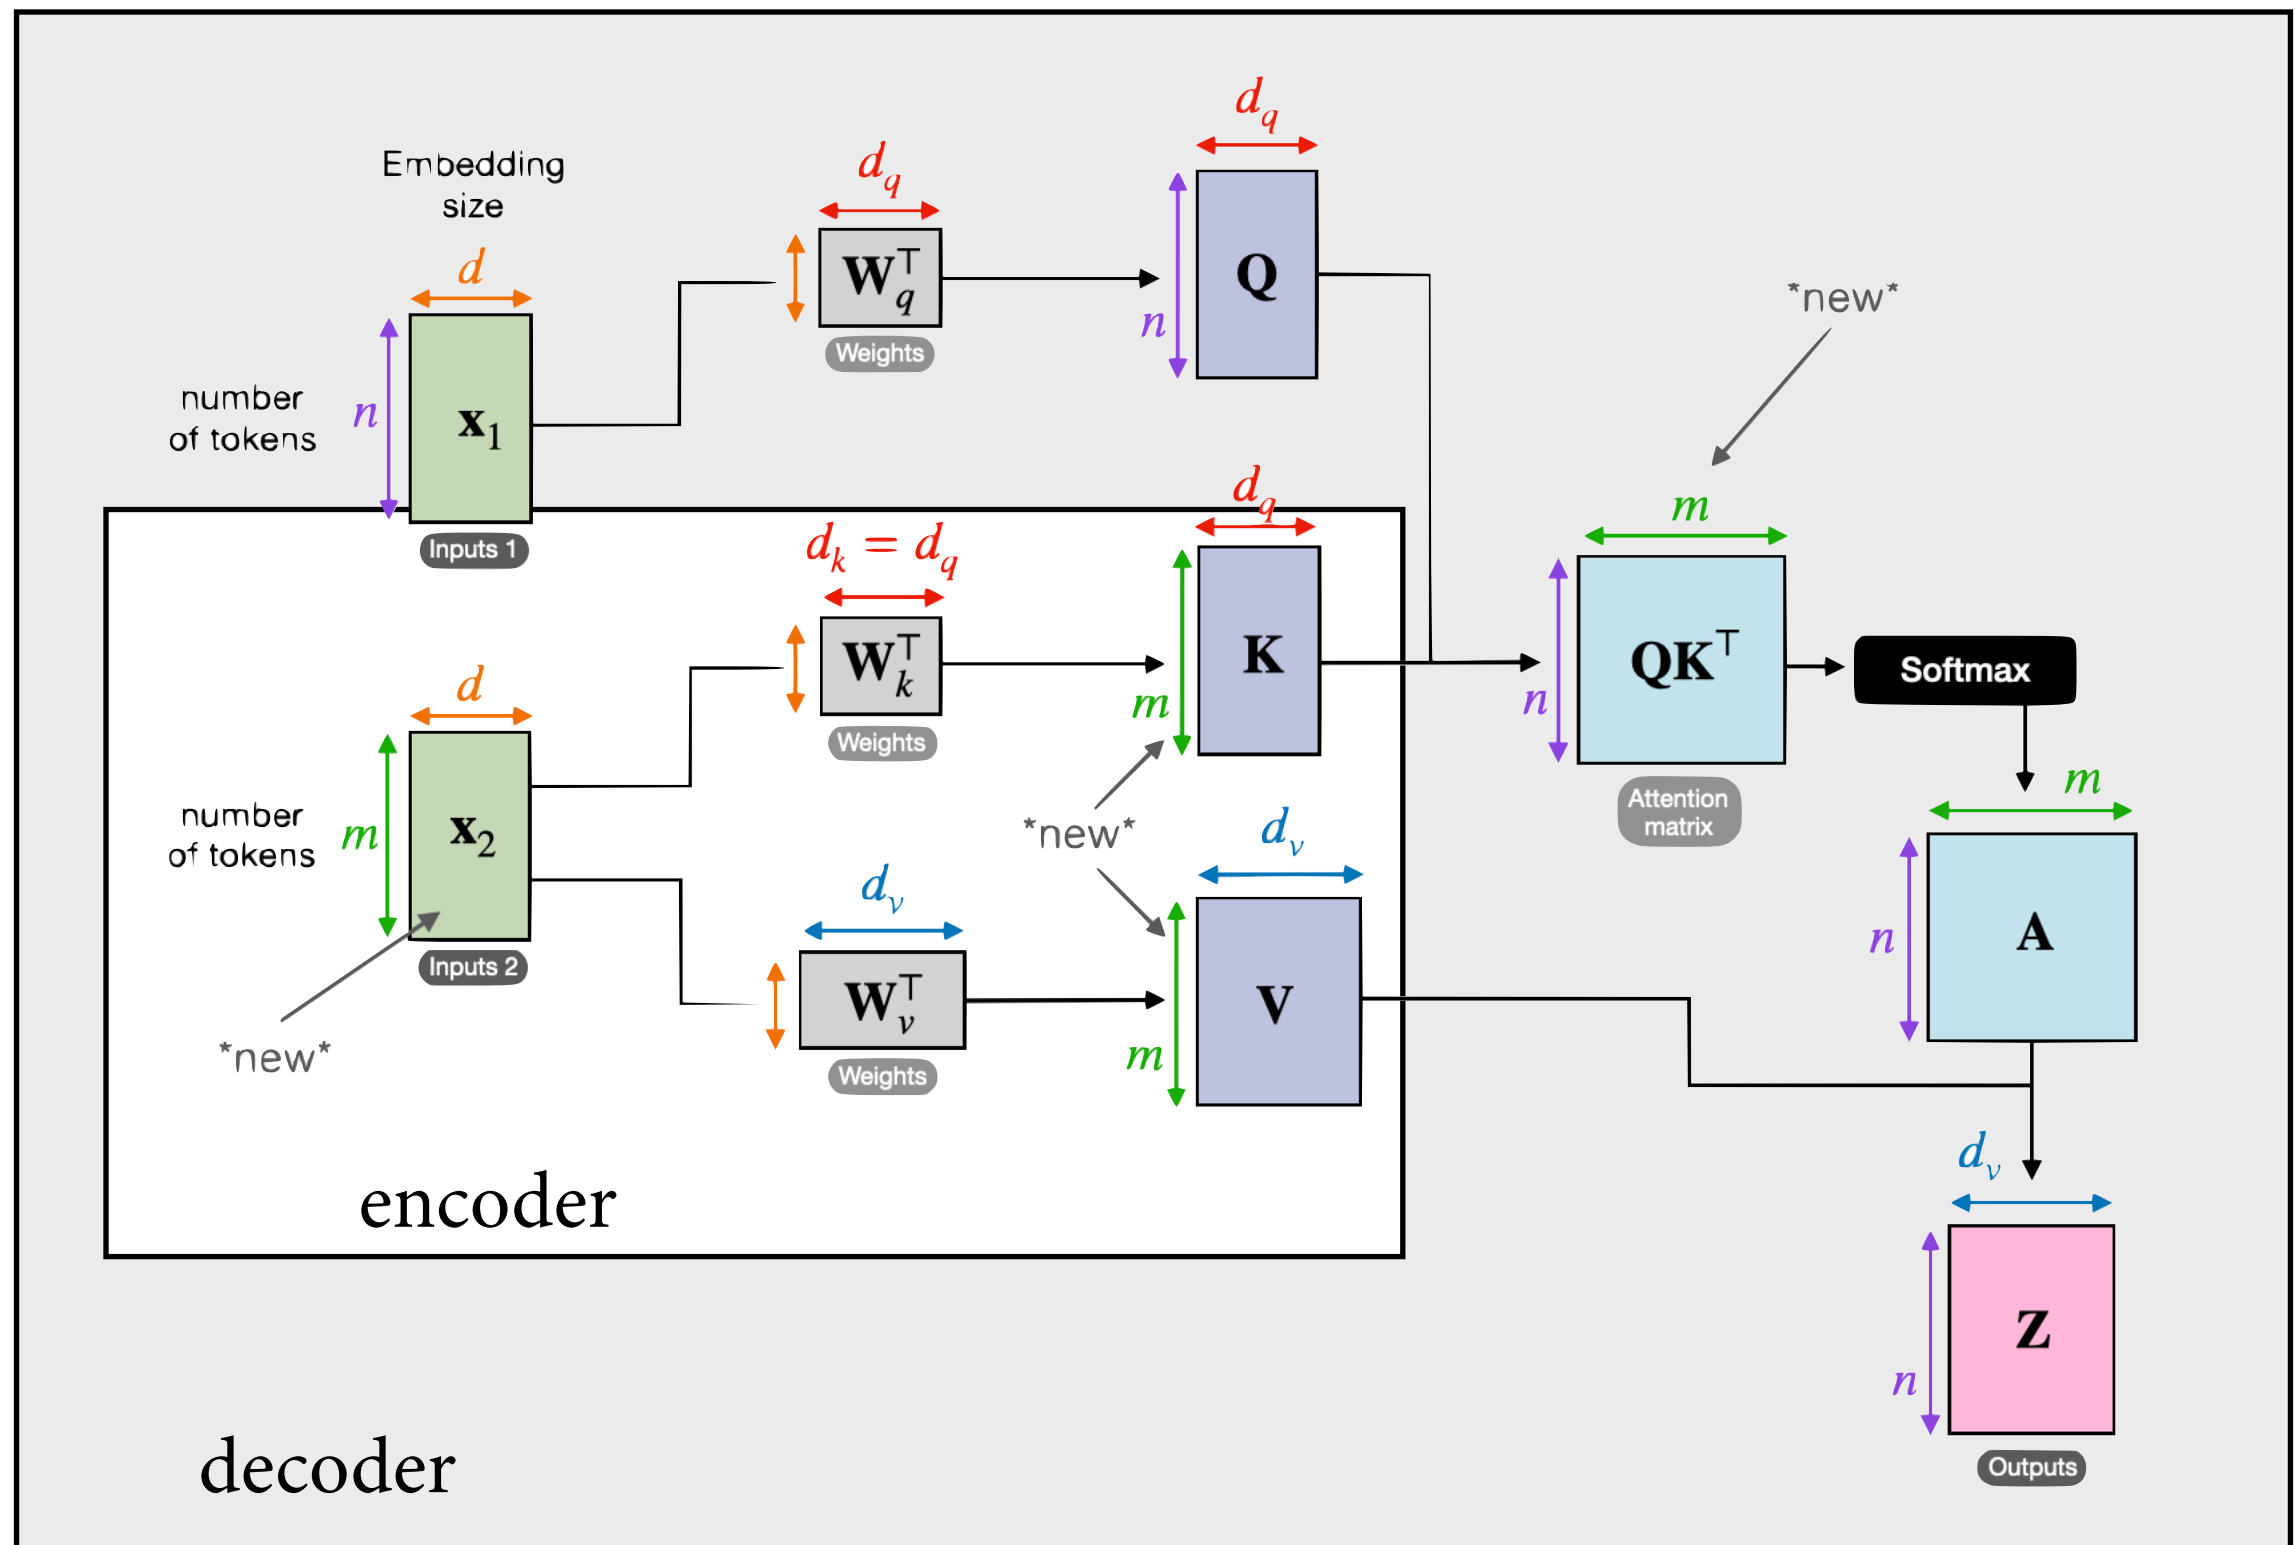
Transformer Decoders



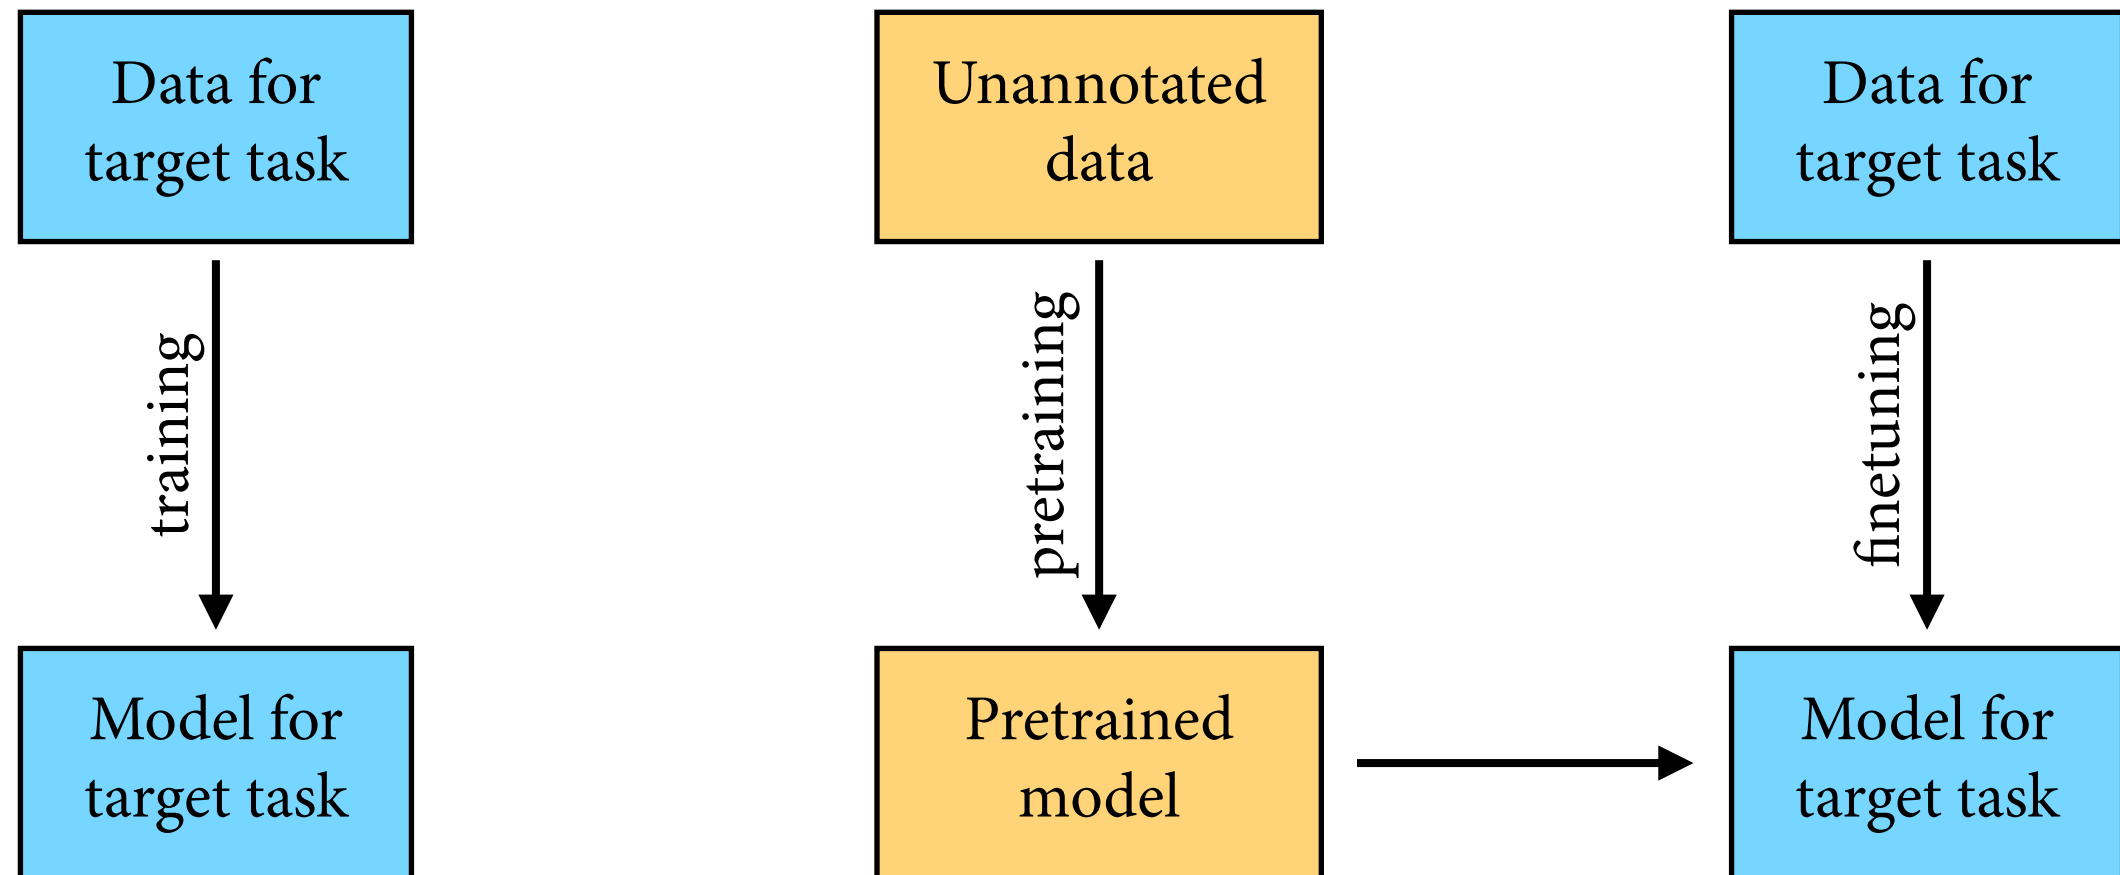
Decoder differences

- Decoder produces outputs token by token.
Each token is appended to the decoder input.
- Decoder can have self-attention to its own inputs - but only to inputs to the left of the current token.
- Decoder also has cross-attention ("encoder-decoder attention") to read encoder outputs.

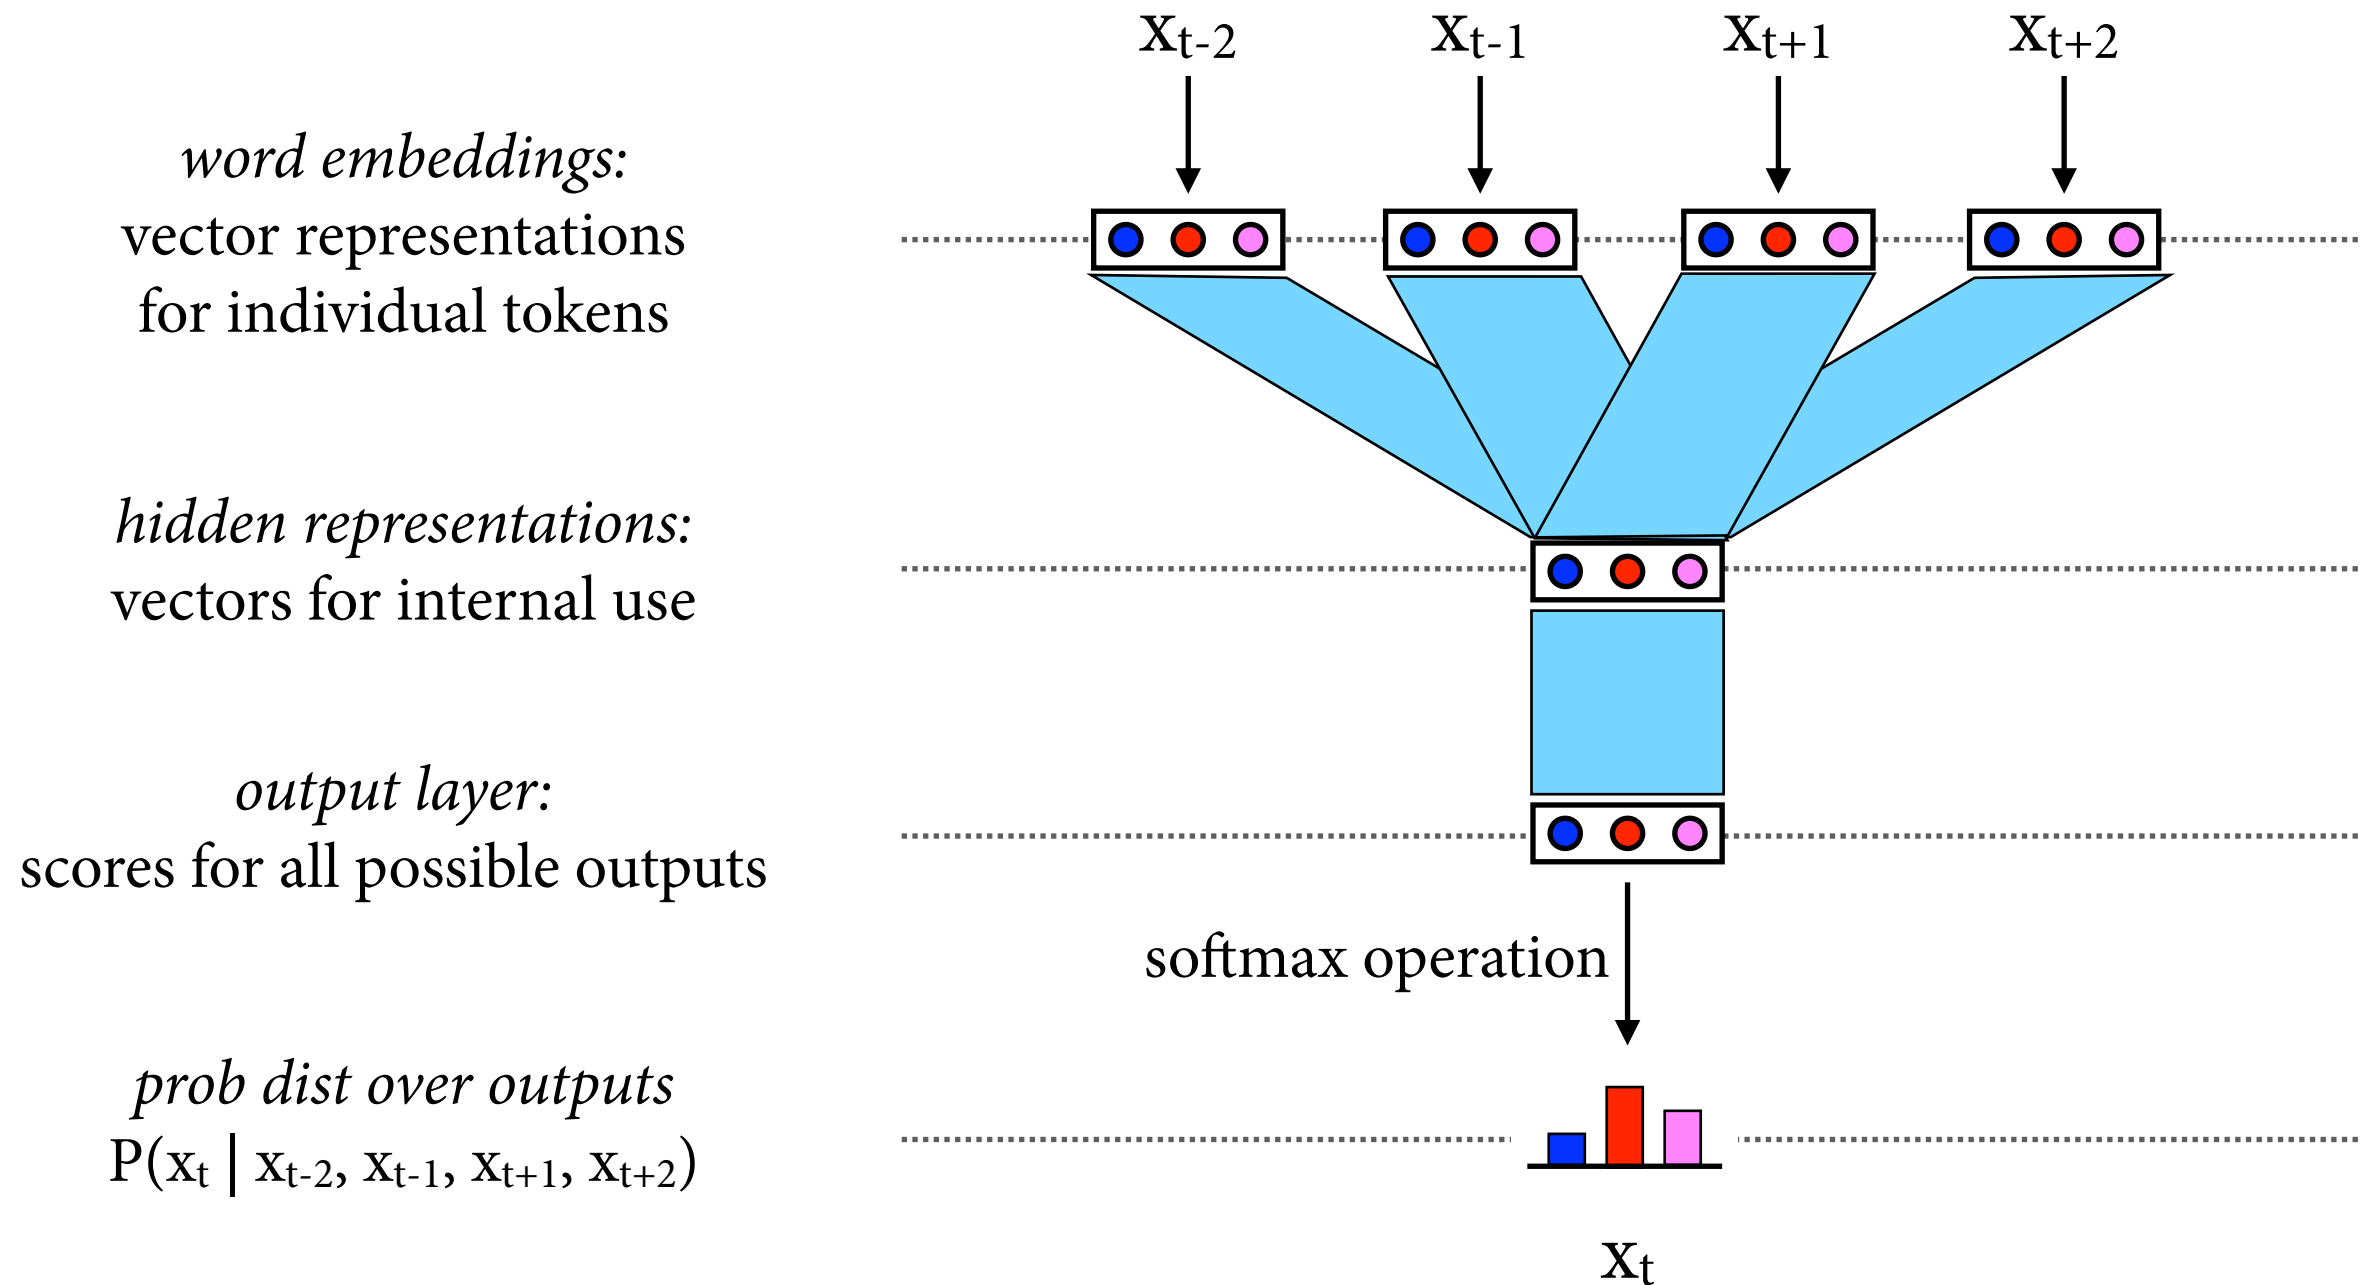
Cross-attention dimensions



Pretraining




Pretrained word embeddings



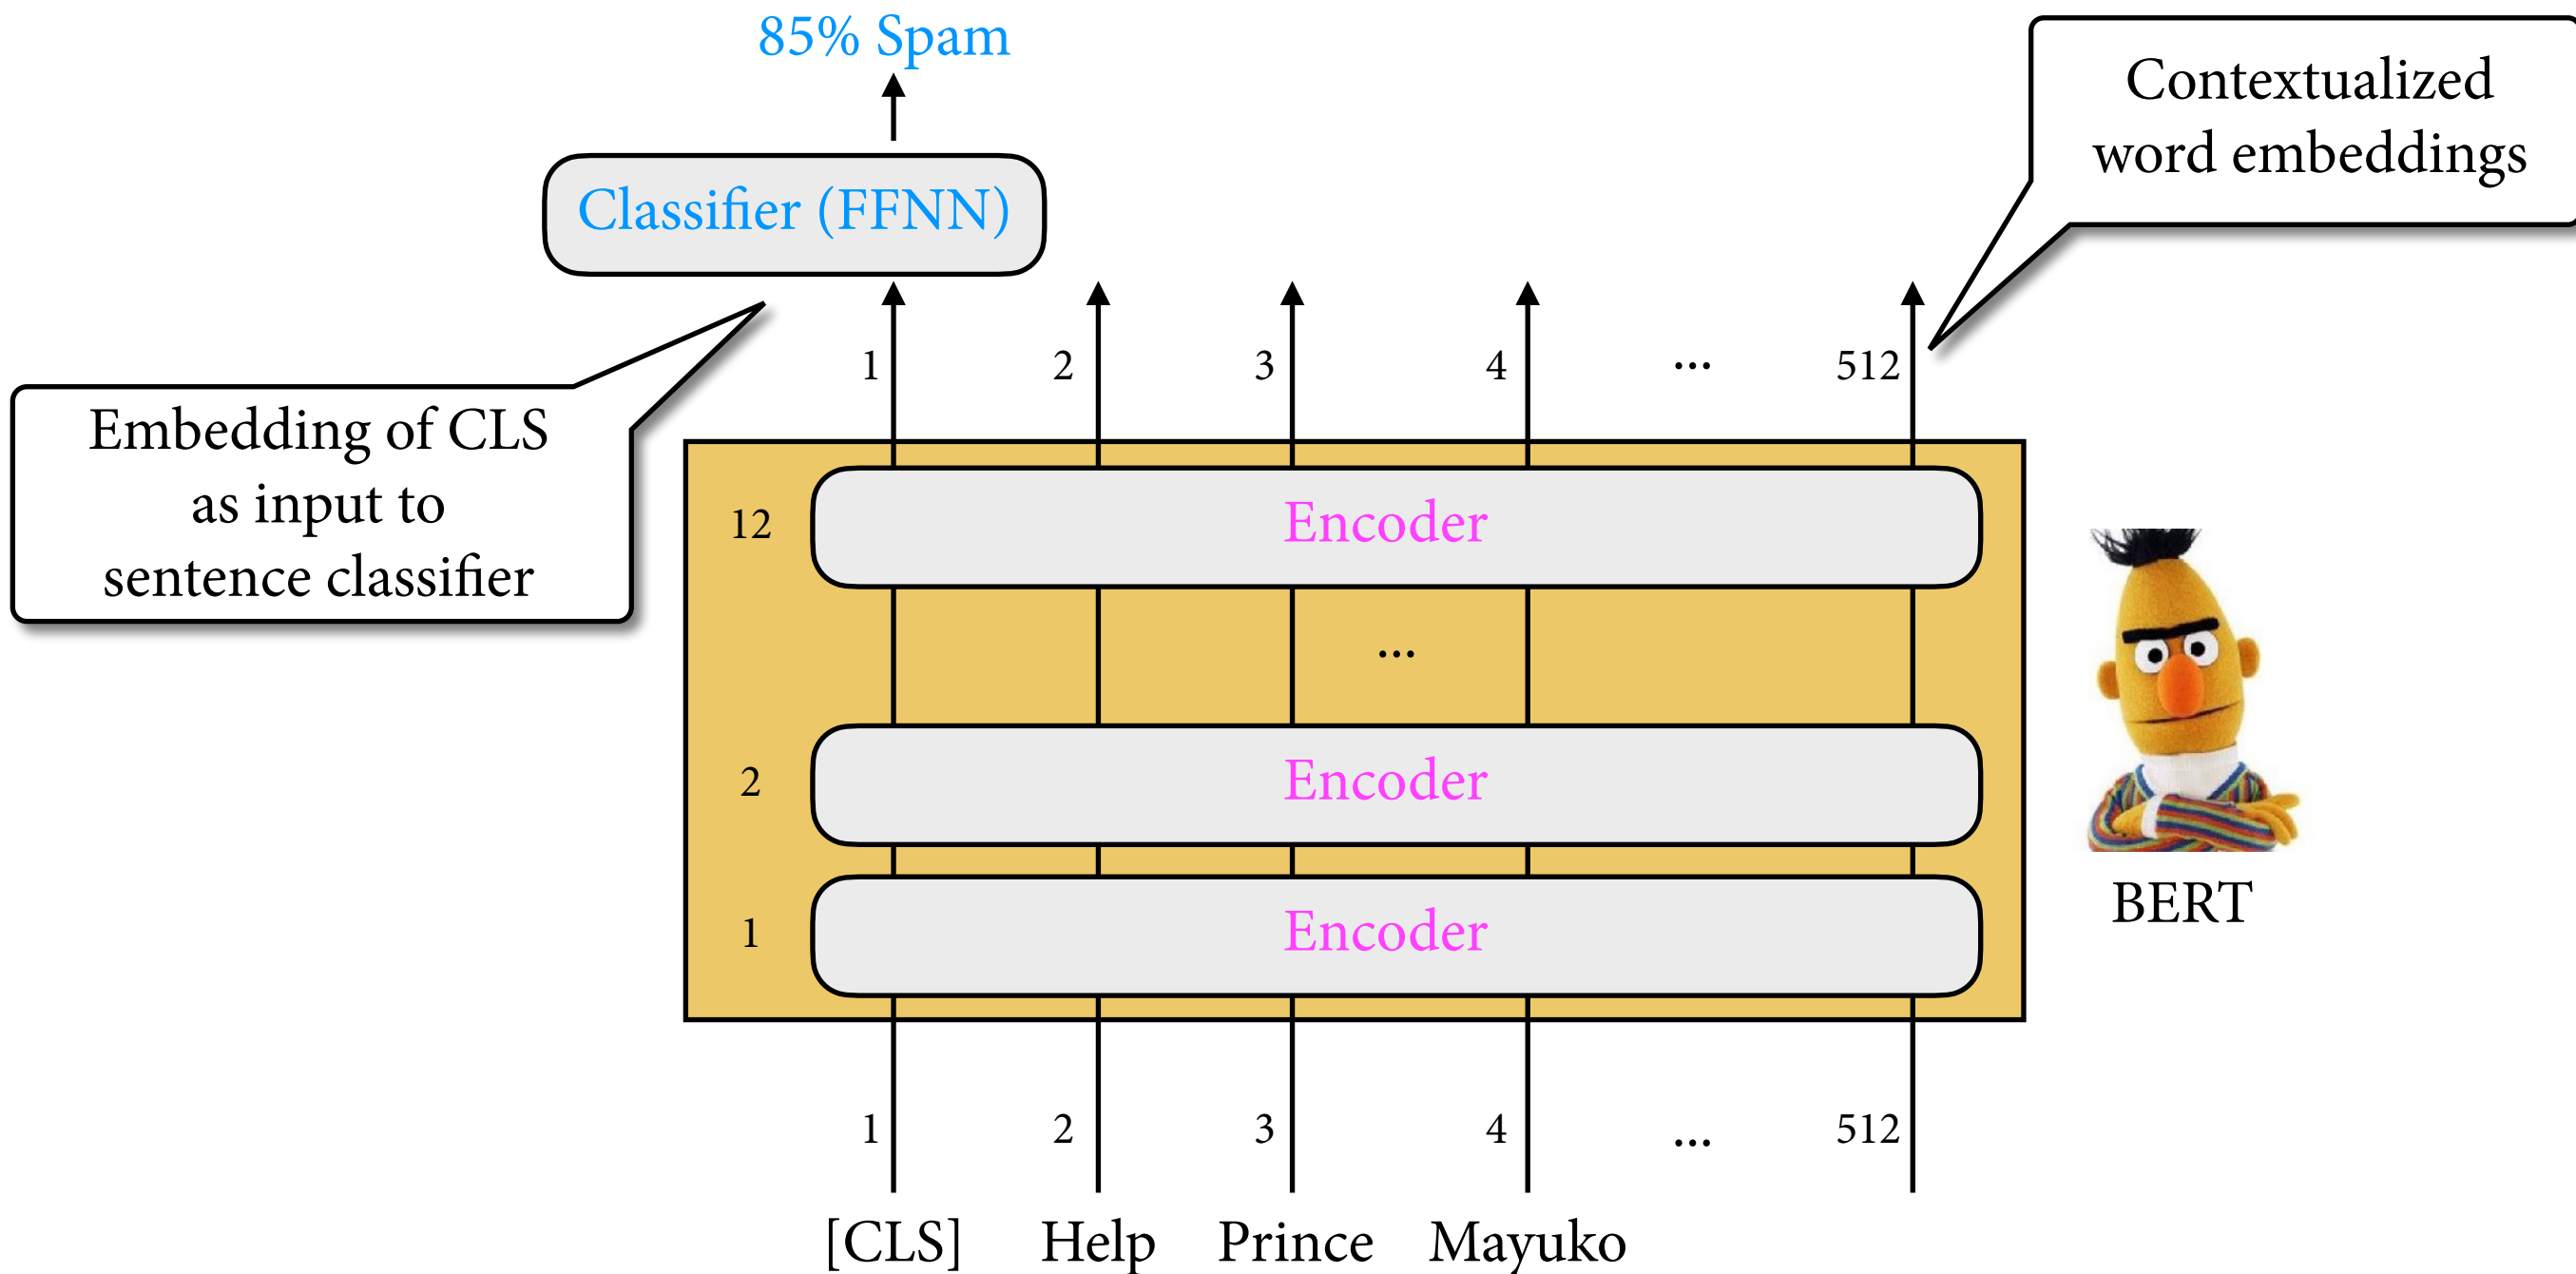
Simplified version of CBOW model (Mikolov et al. 13). Vectors at different layers have different lengths.

Pretrained word embeddings

- Classical word embeddings:
 - ▶ word2vec (2013)
 - ▶ GloVe (2014)
 - ▶ Fasttext (2016)
- All of these map words (or subword tokens) into vector representations.
- Representation depends only on word, not context.



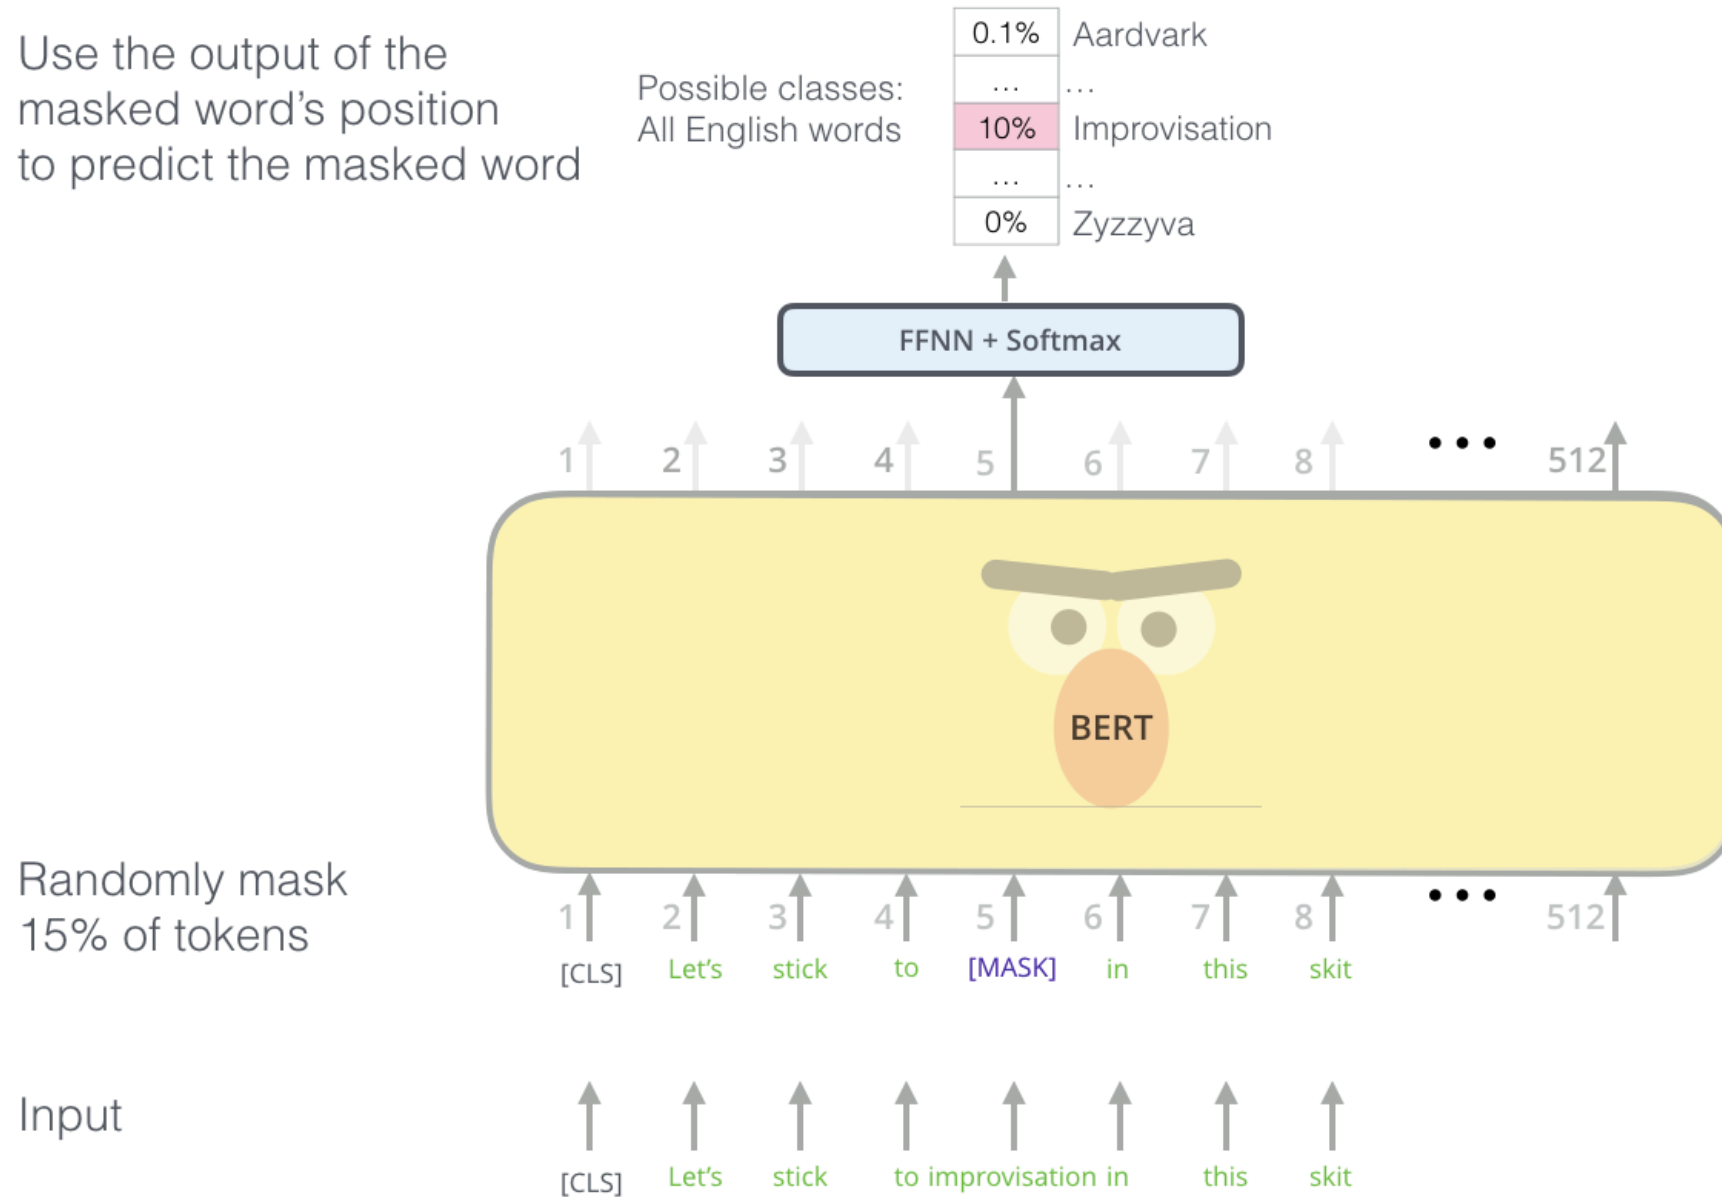
Pretrained transformers



BERT pretraining task

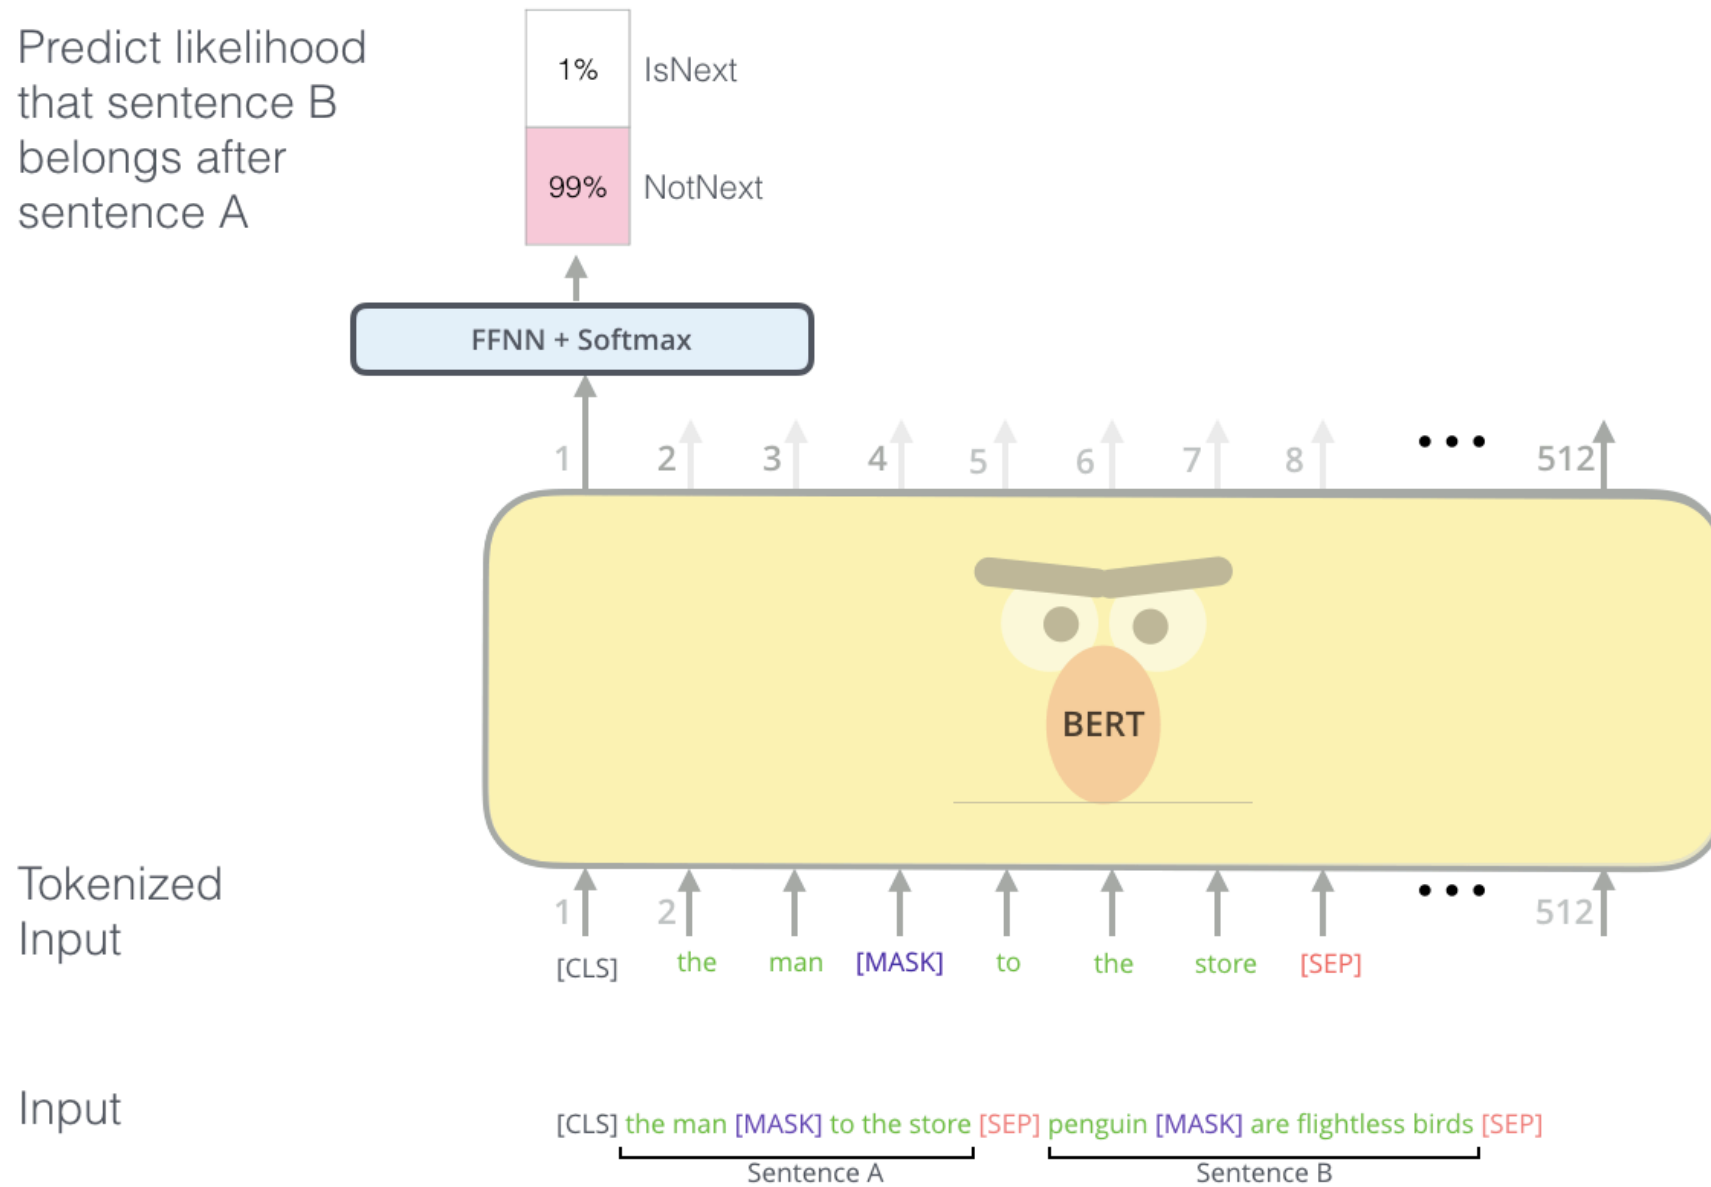
(1) Masked language modeling

Use the output of the masked word's position to predict the masked word



BERT pretraining task

(2) Next-sentence prediction

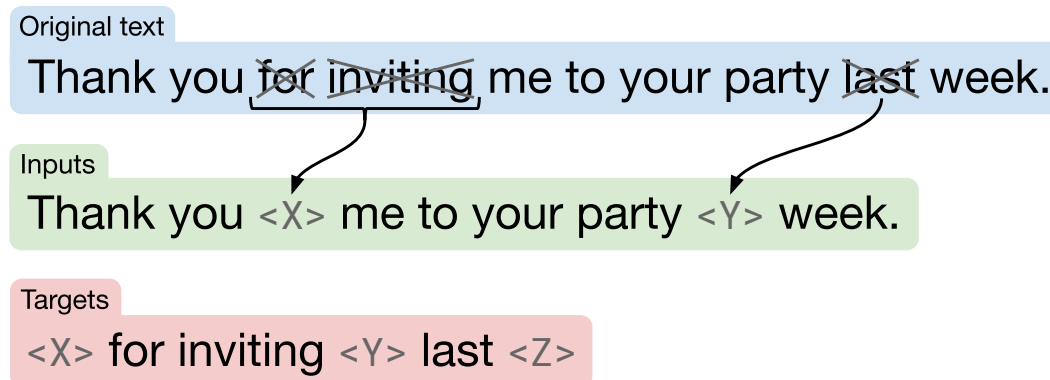


Usefulness of BERT

	DM		PAS		PSD		EDS		AMR 2015	AMR 2017
	id F	ood F	id F	ood F	id F	ood F	Smatch F	EDM	Smatch F	Smatch F
Groschwitz et al. (2018)	-	-	-	-	-	-	-	-	70.2	71.0
Lyu and Titov (2018)	-	-	-	-	-	-	-	-	73.7	74.4 ± 0.16
Zhang et al. (2019)	-	-	-	-	-	-	-	-	-	76.3 ± 0.1
Peng et al. (2017) Basic	89.4	84.5	92.2	88.3	77.6	75.3	-	-	-	-
Dozat and Manning (2018)	93.7	88.9	94.0	90.8	81.0	79.4	-	-	-	-
Buys and Blunsom (2017)	-	-	-	-	-	-	85.5	85.9	60.1	-
Chen et al. (2018)	-	-	-	-	-	-	90.9 ^{1,2}	90.4 ¹	-	-
This paper (GloVe)	90.4 ± 0.2	84.3 ± 0.2	91.4 ± 0.1	86.6 ± 0.1	78.1 ± 0.2	74.5 ± 0.2	87.6 ± 0.1	82.5 ± 0.1	69.2 ± 0.4	70.7 ± 0.2
This paper (BERT)	93.9 ± 0.1	90.3 ± 0.1	94.5 ± 0.1	92.5 ± 0.1	82.0 ± 0.1	81.5 ± 0.3	90.1 ± 0.1	84.9 ± 0.1	74.3 ± 0.2	75.3 ± 0.2
Peng et al. (2017) Freda1	90.0	84.9	92.3	88.3	78.1	75.8	-	-	-	-
Peng et al. (2017) Freda3	90.4	85.3	92.7	89.0	78.5	76.4	-	-	-	-
This paper, MTL (GloVe)	91.2 ± 0.1	85.7 ± 0.0	92.2 ± 0.2	88.0 ± 0.3	78.9 ± 0.3	76.2 ± 0.4	88.2 ± 0.1	83.3 ± 0.1	(70.4) ³ ± 0.2	71.2 ± 0.2
This paper, MTL (BERT)	94.1 ± 0.1	90.5 ± 0.1	94.7 ± 0.1	92.8 ± 0.1	82.1 ± 0.2	81.6 ± 0.1	90.4 ± 0.1	85.2 ± 0.1	(74.5) ³ ± 0.1	75.3 ± 0.1

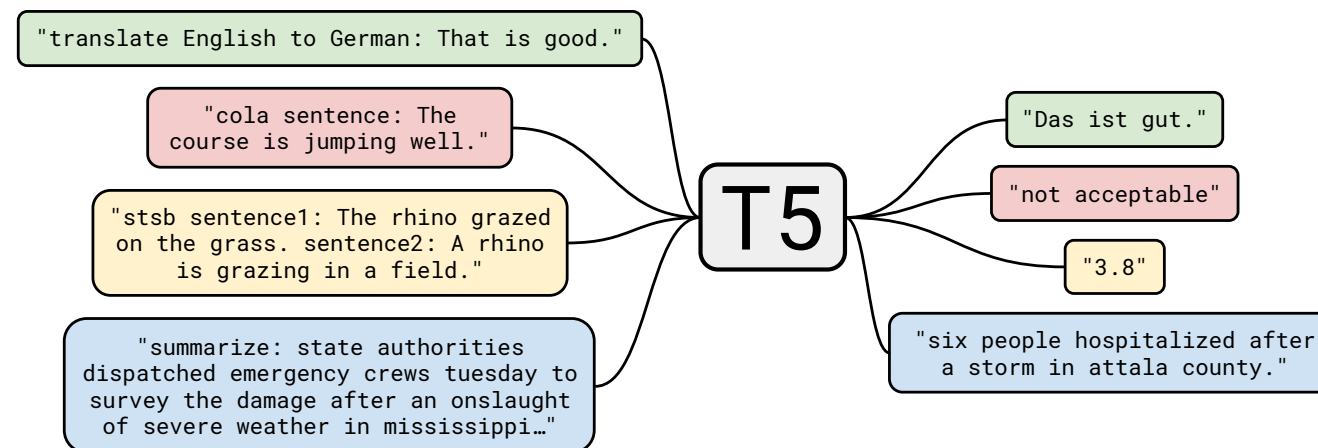
Pretrained enc-dec models

- T5: Standard encoder-decoder transformer.
- Unsupervised pretraining on C4 corpus: cleaned-up version of the Common Crawl.
- Training objective: Denoising.



Finetuning T5

- Map wide variety of NLP tasks into seq2seq task:

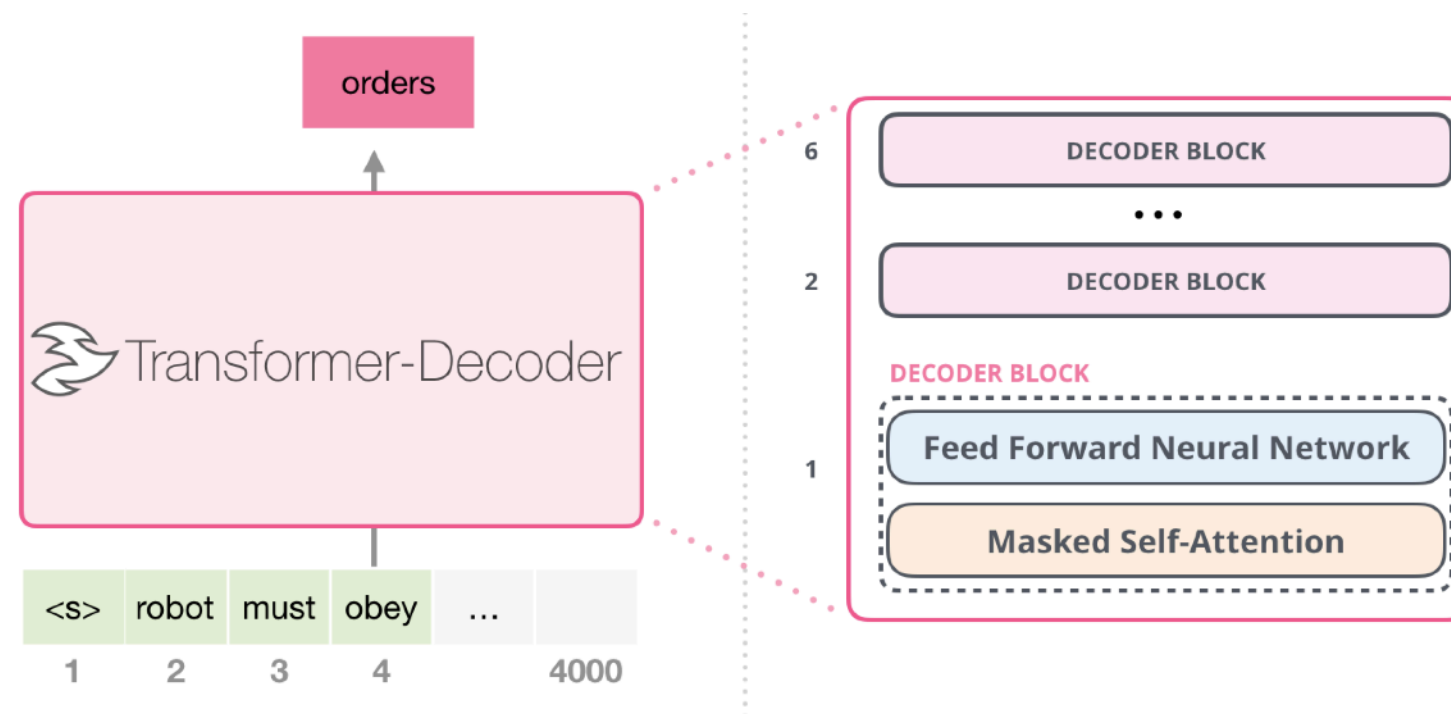


- Apply T5 to new task by finetuning all of its parameters:

Fine-tuning method	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ All parameters	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Adapter layers, $d = 32$	80.52	15.08	79.32	60.40	13.84	17.88	15.54
Adapter layers, $d = 128$	81.51	16.62	79.47	63.03	19.83	27.50	22.63
Adapter layers, $d = 512$	81.54	17.78	79.18	64.30	23.45	33.98	25.81
Adapter layers, $d = 2048$	81.51	16.62	79.47	63.03	19.83	27.50	22.63
Gradual unfreezing	82.50	18.95	79.17	70.79	26.71	39.02	26.93

Decoder-only models

- GPT models: Use only a stack of decoders (GPT-2 XL has 48 of them).



- No cross-attention because there is no encoder.
- *Masked* self-attention: can only attend to the left.

In-context learning

The three settings we explore for in-context learning

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```

Traditional fine-tuning (not used for GPT-3)

Fine-tuning

The model is trained via repeated gradient updates using a large corpus of example tasks.



(Brown et al. 2020: GPT-3 as "few-shot learner")

Tokenization

- Challenge in broad-coverage neural LMs:
How do you keep vocabulary size under control?
- Standard solution nowadays is to use subword tokenizers, which break words up into reusable pieces.
 - ▶ Byte-pair encoding (BPE, Sennrich et al. 2016)
 - ▶ WordPiece (introduced for BERT, Devlin et al. 2019)
 - ▶ SentencePiece is a popular implementation
 - ▶ Huggingface models such as XML-RoBERTa come with their own subword tokenizers already implemented.

Byte-pair encoding

Original corpus:

hug	pug	pun	bun	hugs
x 10	x 5	x 12	x 4	x 5

Split into characters:

h u g	p u g	p u n	b u n	h u g s
x 10	x 5	x 12	x 4	x 5

Merge most frequent token pair:

h ug	p ug	p u n	b u n	h ug s
x 10	x 5	x 12	x 4	x 5

(repeatedly)

h ug	p ug	p un	b un	h ug s
x 10	x 5	x 12	x 4	x 5

Stop at intended vocabulary size:

hug	p ug	p un	b un	hug s
x 10	x 5	x 12	x 4	x 5

Summary

- Transformers: sequence models that can be trained in parallel.
- Pretraining is extremely effective method, especially for very large transformers.
 - ▶ finetuning vs. in-context learning
- Tokenization is a challenge which is often addressed with subword tokenization.