

Semantic parsing

Computational Linguistics


Alexander Koller

09 January 2024

Overview

1. Semantics in computational linguistics (up to 2000).
2. Classical semantic parsing (Geoquery, early 2000s).
3. Semantic graph parsing:
 - graph-based (JAMR, 2014)
 - compositional (AM dependency parsing, 2019)
 - finetuning of pretrained models (SPRING, 2021)

Computing with meanings



- Ancient problem: *inference*.
 - ▶ How can we tell whether a sentence follows from others?
 - ▶ Can we compute this automatically?

All men are mortal.

Socrates is a man.

Therefore, Socrates is mortal.


(Syllogism, Aristotle, 350 BC)

At the other end of Pennsylvania Avenue,
people began to line up for a White House tour.

People formed a line at the end of
Pennsylvania Avenue.

(MNLI dataset, 2018)

Formal meaning representations



- Aristotle with more modern tools (ca. 2000):
 - ▶ Compute *meaning representation* in some formal language (e.g. predicate logic)
 - ▶ so that it captures something relevant about the sentence's meaning (e.g. its *truth conditions*)
 - ▶ and then use reasoning tools for the formal language (e.g. a *theorem prover* for predicate logic)

All men are mortal.

Socrates is a man.


Therefore, Socrates is mortal.

$\forall x. \text{man}(x) \rightarrow \text{mortal}(x)$

$\text{man}(s)$

$\text{mortal}(s)$

Compositional semantics



$S \rightarrow NP VP$

$\langle S \rangle = \langle NP \rangle (\langle VP \rangle)$

$VP \rightarrow V NP$

$\langle VP \rangle = \lambda y \langle NP \rangle (\langle V \rangle (y))$

$NP \rightarrow Det N$

$\langle NP \rangle = \langle Det \rangle (\langle N \rangle)$

$NP \rightarrow John$

$\langle NP \rangle = \lambda P P(j^*)$

$V \rightarrow eats$

$\langle V \rangle = eat'$

$Det \rightarrow a$

$\langle Det \rangle = \lambda P \lambda Q \exists x P(x) \wedge Q(x)$

$N \rightarrow sandwich$

$\langle N \rangle = sw'$

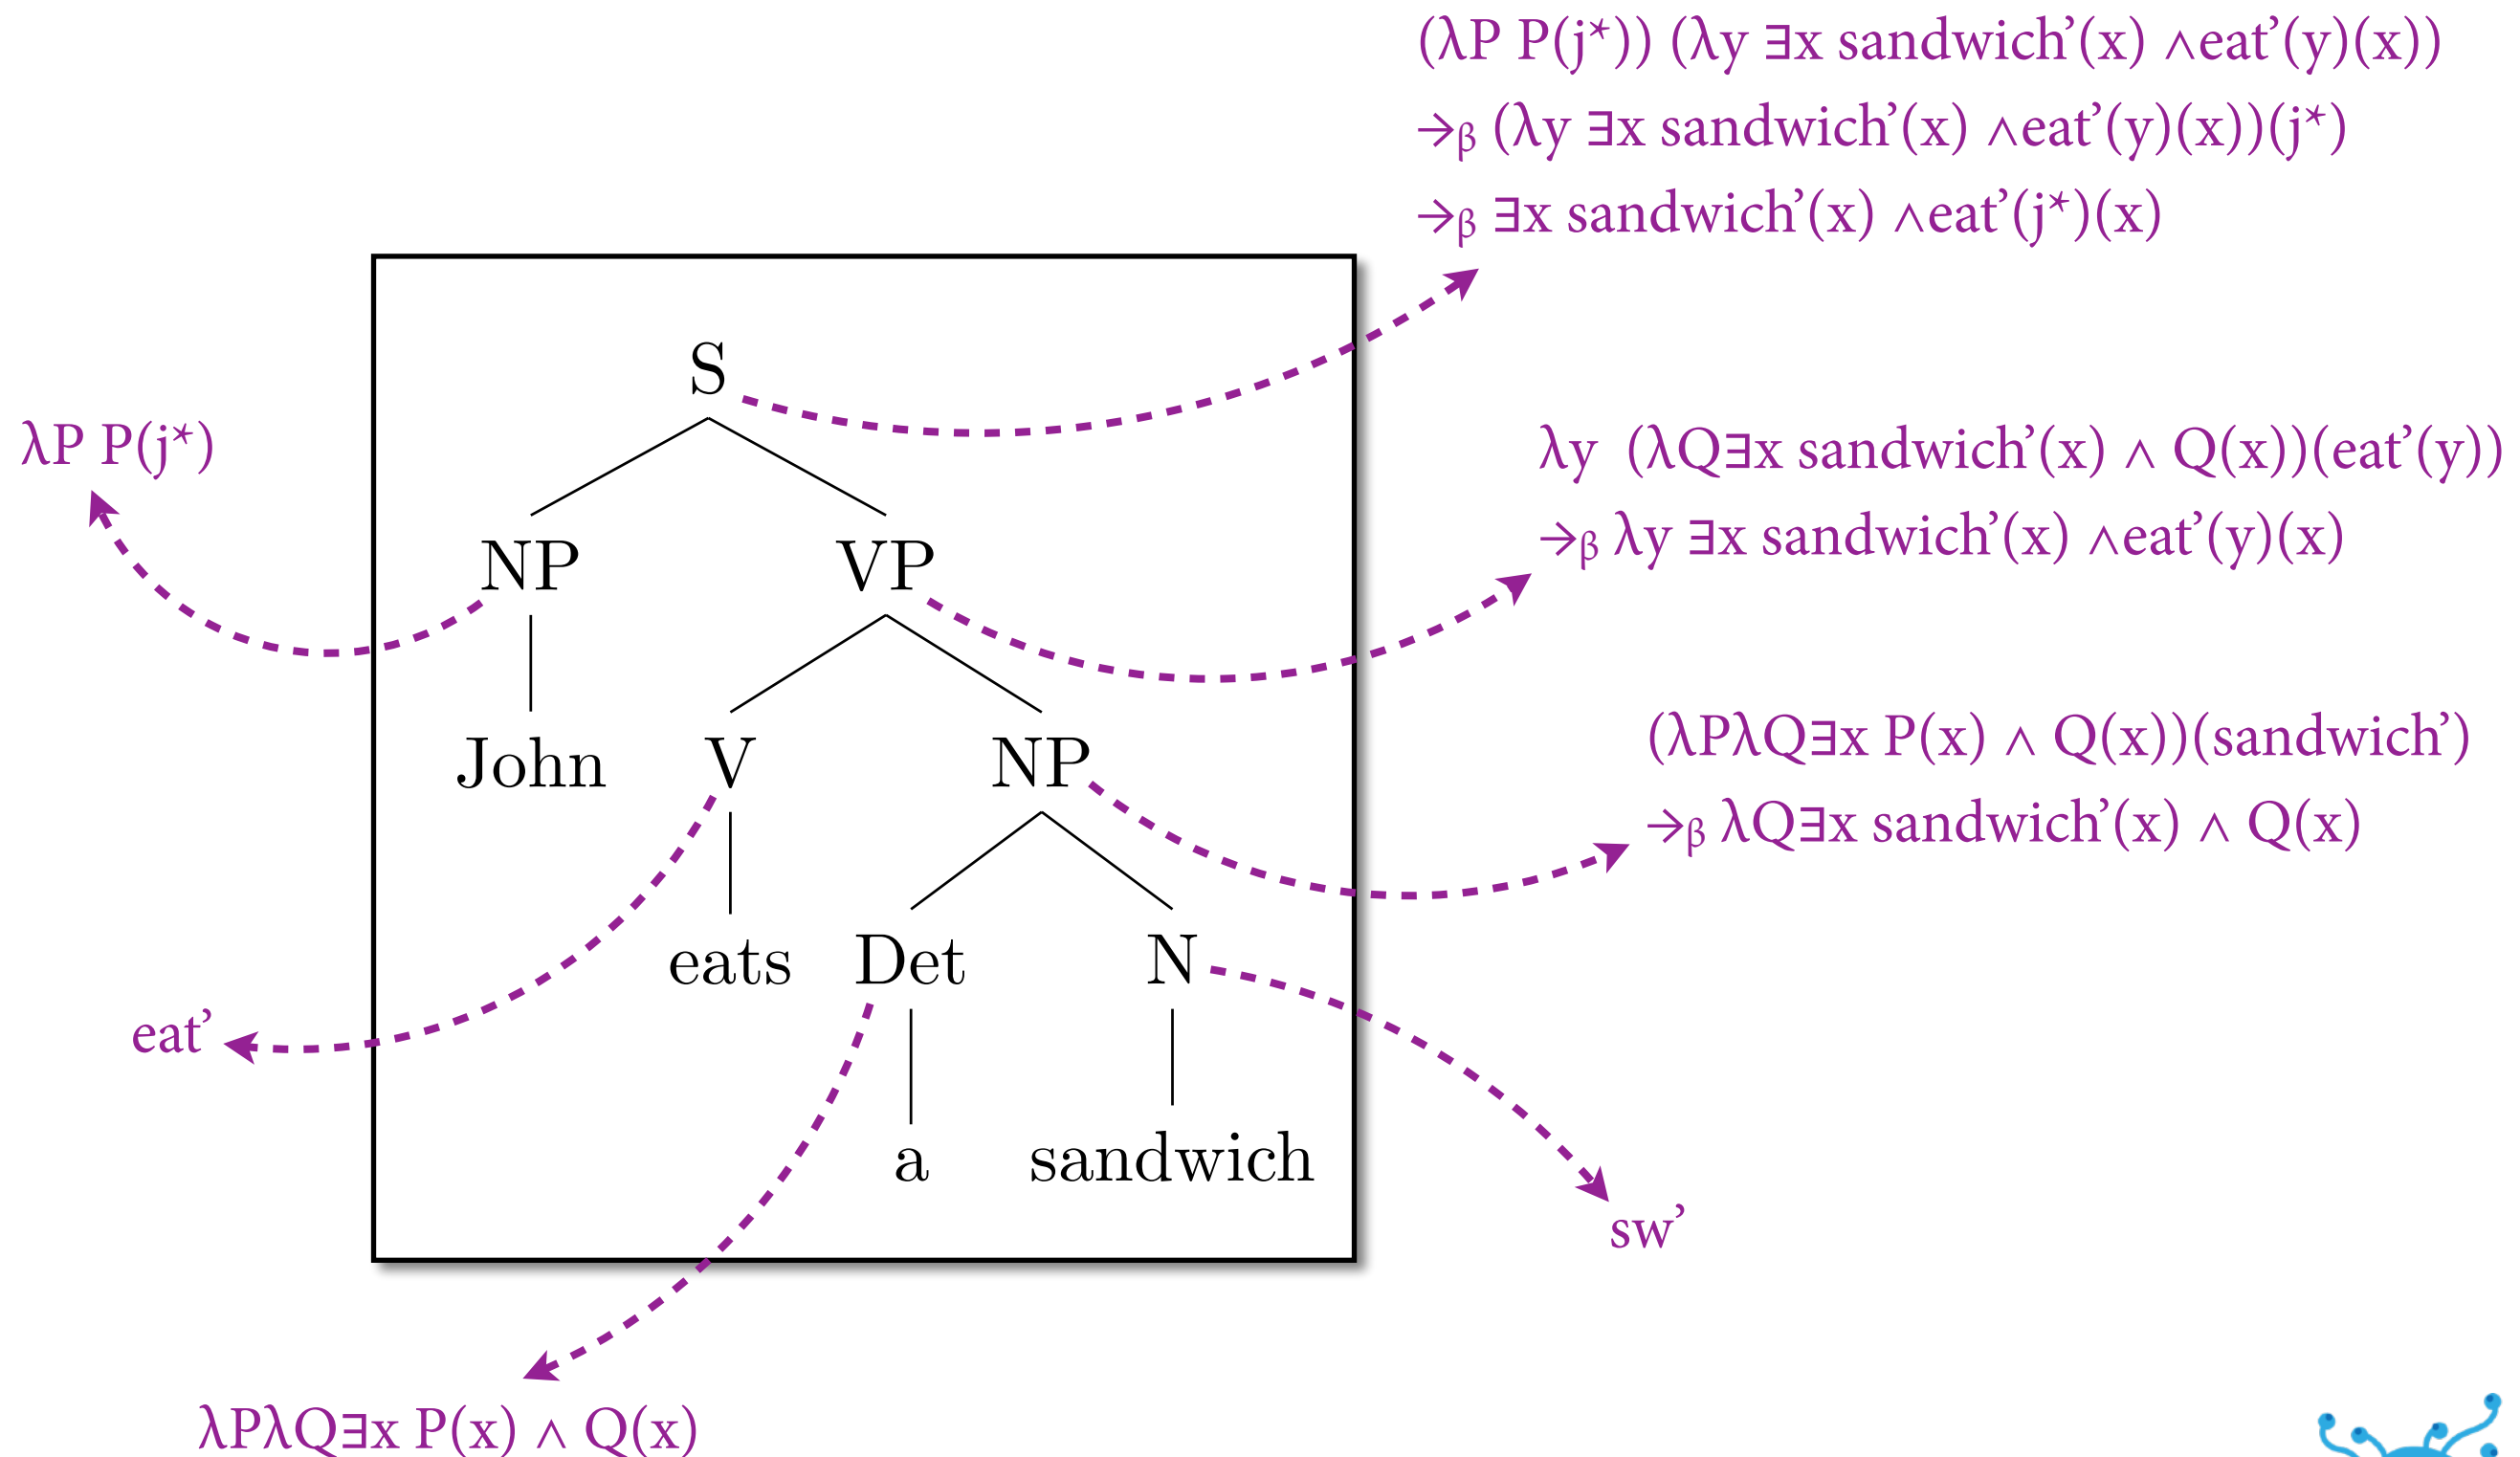


when you apply this
syntax rule ...



... construct λ -term for parent
from λ -terms for children like this

Example



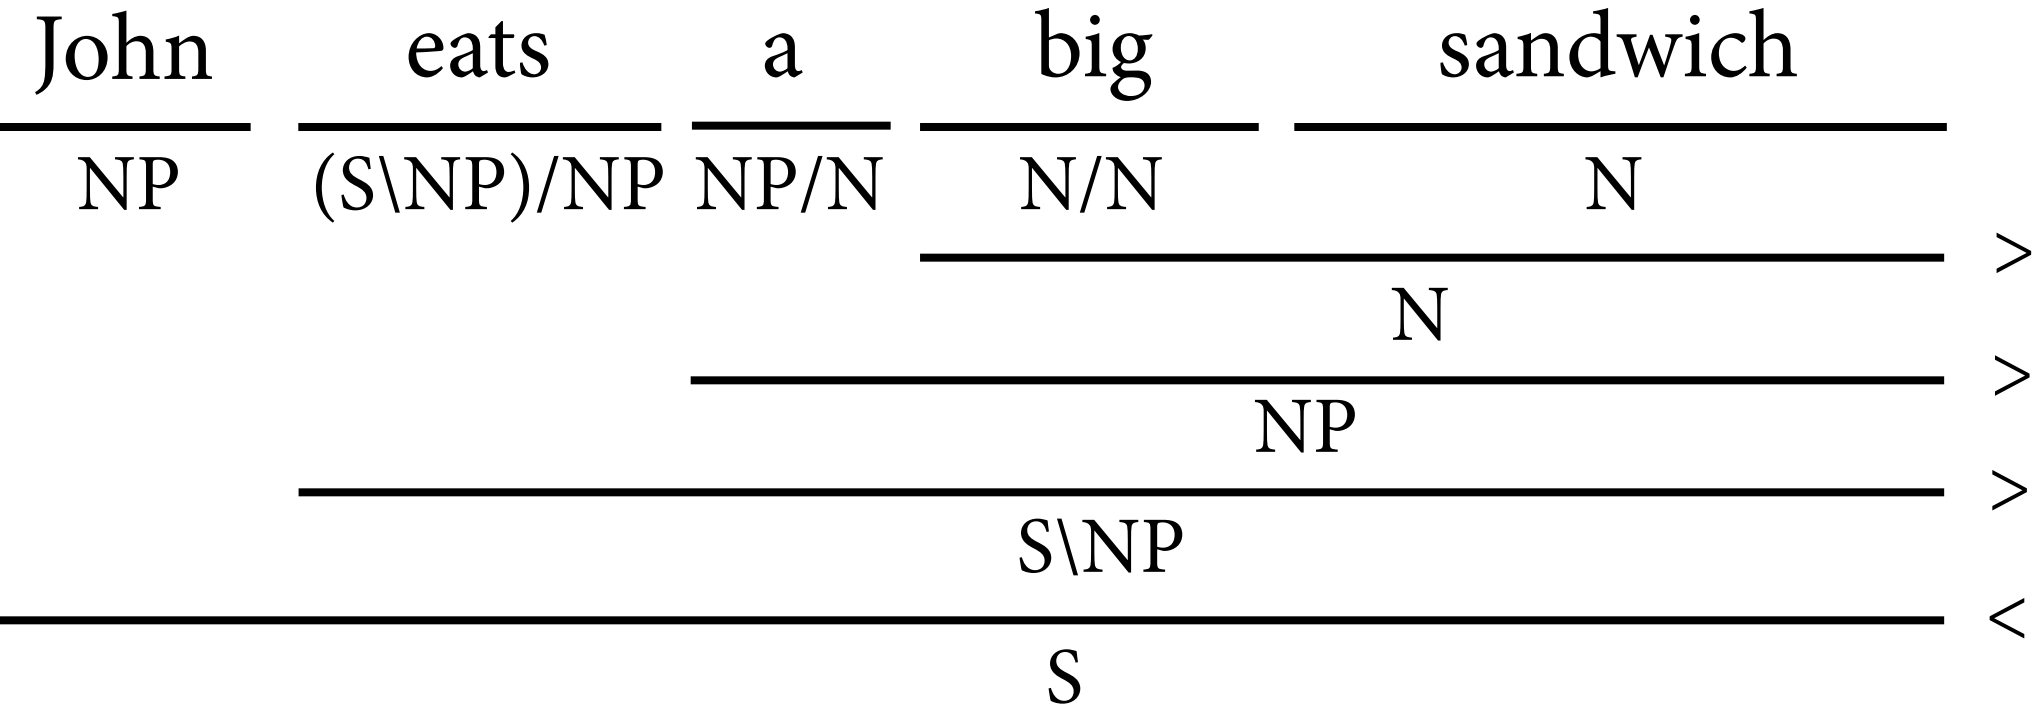
Semantic parsing

- Open issue in classical semantics construction:
Where do we get large grammar that supports it?
- Current approach in CL is *semantic parsing*:
learn mapping from sentence to formal meaning
representation using statistical methods.
- Started with Geoquery corpus (880 sentences):

What is the smallest state by area?

```
answer(x1, smallest(x2, state(x1), area(x1, x2)))
```

Combinatory categorial grammar



Semantics in CCG

$\frac{X: a}{Y/(Y \backslash X): \lambda P.P(a)} >T$	$\frac{X/Y: f \quad Y/Z: g}{X/Z: \lambda x.f(g(x))} >B$	$\frac{X/Y: f \quad Y \backslash Z: g}{X \backslash Z: \lambda x.f(g(x))} >Bx$
$\frac{X: a}{Y \backslash (Y/X): \lambda P.P(a)} <T$	$\frac{Y \backslash Z: g \quad X \backslash Y: f}{X \backslash Z: \lambda x.f(g(x))} <B$	$\frac{Y/Z: g \quad X \backslash Y: f}{X/Z: \lambda x.f(g(x))} <Bx$

$\frac{\text{John}}{\text{NP: } h^*}$	$\frac{\text{eats}}{(\text{S} \backslash \text{NP})/\text{NP: } eat'}$	$\frac{\text{a sandwich}}{\text{NP: } sandwich'}$
$\frac{\text{S}/(\text{S} \backslash \text{NP}): \lambda P.P(h^*)}{\text{S}/\text{NP: } \lambda x.(\lambda P.P(h^*))(eat'(x)) \Rightarrow_{\beta} \lambda x.eat'(x)(h^*)} >T$	$>B$	
$\frac{\text{S: } (\lambda x.eat'(x)(h^*))(sandwich') \Rightarrow_{\beta} eat'(sandwich')(h^*)}{\text{S: } (\lambda x.eat'(x)(h^*))(sandwich') \Rightarrow_{\beta} eat'(sandwich')(h^*)} >$		

Zettlemoyer & Collins

GENLEX: build candidates for lexicon entries



Rules		Categories produced from logical form
Input Trigger	Output Category	$\arg \max(\lambda x.state(x) \wedge borders(x, texas), \lambda x.size(x))$
constant c	$NP : c$	$NP : texas$
arity one predicate p_1	$N : \lambda x.p_1(x)$	$N : \lambda x.state(x)$
arity one predicate p_1	$S \backslash NP : \lambda x.p_1(x)$	$S \backslash NP : \lambda x.state(x)$
arity two predicate p_2	$(S \backslash NP) / NP : \lambda x.\lambda y.p_2(y, x)$	$(S \backslash NP) / NP : \lambda x.\lambda y.borders(y, x)$
arity two predicate p_2	$(S \backslash NP) / NP : \lambda x.\lambda y.p_2(x, y)$	$(S \backslash NP) / NP : \lambda x.\lambda y.borders(x, y)$
arity one predicate p_1	$N / N : \lambda g.\lambda x.p_1(x) \wedge g(x)$	$N / N : \lambda g.\lambda x.state(x) \wedge g(x)$
literal with arity two predicate p_2 and constant second argument c	$N / N : \lambda g.\lambda x.p_2(x, c) \wedge g(x)$	$N / N : \lambda g.\lambda x.borders(x, texas) \wedge g(x)$
arity two predicate p_2	$(N \backslash N) / NP : \lambda x.\lambda g.\lambda y.p_2(x, y) \wedge g(x)$	$(N \backslash N) / NP : \lambda g.\lambda x.\lambda y.borders(x, y) \wedge g(x)$
an $\arg \max$ / \min with second argument arity one function f	$NP / N : \lambda g.\arg \max / \min(g, \lambda x.f(x))$	$NP / N : \lambda g.\arg \max(g, \lambda x.size(x))$
an arity one numeric-ranged function f	$S / NP : \lambda x.f(x)$	$S / NP : \lambda x.size(x)$

(Zettlemoyer & Collins 2005, 2007)

Zettlemoyer & Collins

overall learning algorithm

Algorithm:

- For $t = 1 \dots T$

Step 1: (Lexical generation)

- For $i = 1 \dots n$:
 - Set $\lambda = \Lambda_0 \cup \text{GENLEX}(S_i, L_i)$.
 - Calculate $\pi = \text{PARSE}(S_i, L_i, \lambda, \bar{\theta}^{t-1})$.
 - Define λ_i to be the set of lexical entries in π .
- Set $\Lambda_t = \Lambda_0 \cup \bigcup_{i=1}^n \lambda_i$

analogous to Viterbi EM



Step 2: (Parameter Estimation)

- Set $\bar{\theta}^t = \text{ESTIMATE}(\Lambda_t, E, \bar{\theta}^{t-1})$


Λ are CCG lexicons; λ contains new lexicon entries.

π is a CCG parse tree.

θ are parameters of a log-linear probability model.


Evaluation results

System	Variable Free			Lambda Calculus		
	Rec.	Pre.	F1	Rec.	Pre.	F1
Cross Validation Results						
KRISP	71.7	93.3	81.1	—	—	—
WASP	74.8	87.2	80.5	—	—	—
Lu08	81.5	89.3	85.2	—	—	—
λ -WASP	—	—	—	86.6	92.0	89.2
Independent Test Set						
ZC05	—	—	—	79.3	96.3	87.0
ZC07	—	—	—	86.1	91.6	88.8
UBL	81.4	89.4	85.2	85.0	94.1	89.3
UBL-s	84.3	85.2	84.7	87.9	88.5	88.2



Abstract Meaning Representations

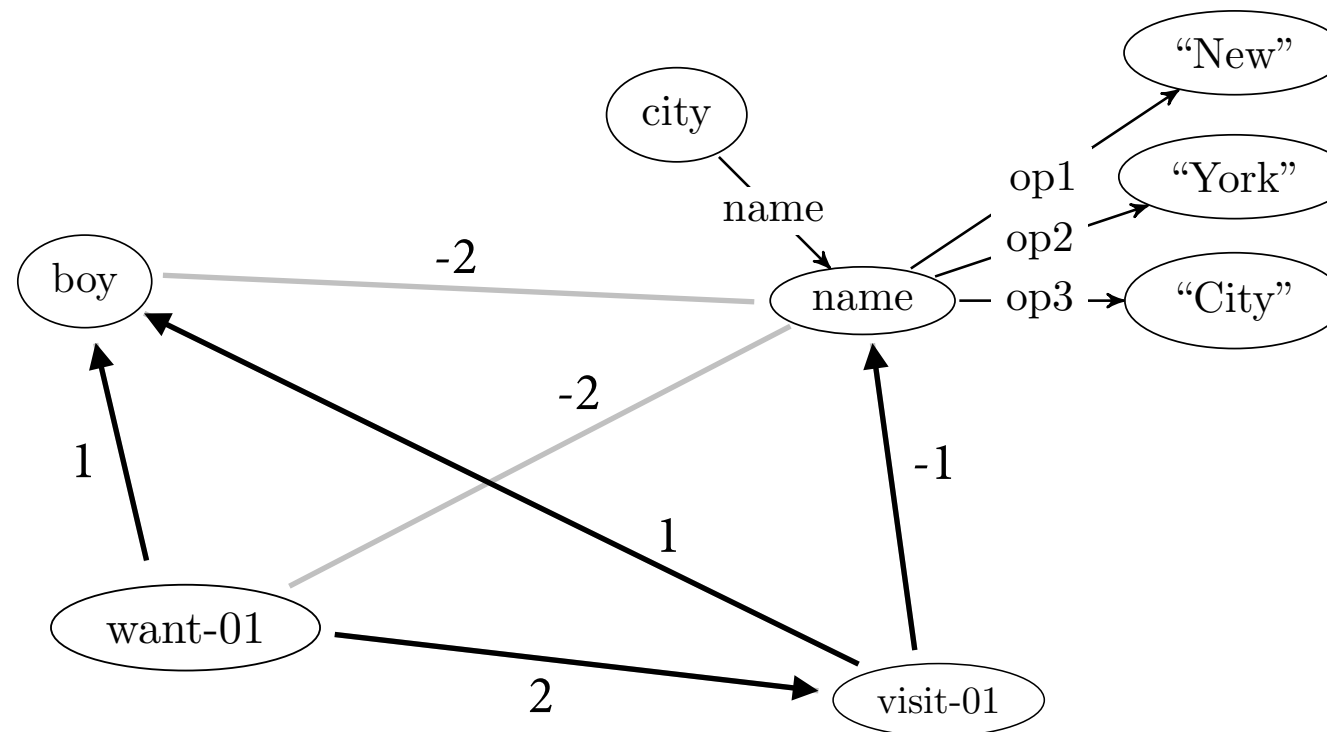
- Limitations of Geoquery:
 - ▶ semantic representations are trees — (too) easy
 - ▶ very small
- Since 2014, much larger corpora available, e.g. AMRBank: ~40k AMRs, graphs as semantic representations.



“I don’t want anyone to read my book carelessly.”

Dependency-style AMR parsing

“The boy wants to visit New York City.”



Concept Identification: determine atomic graph for each word.

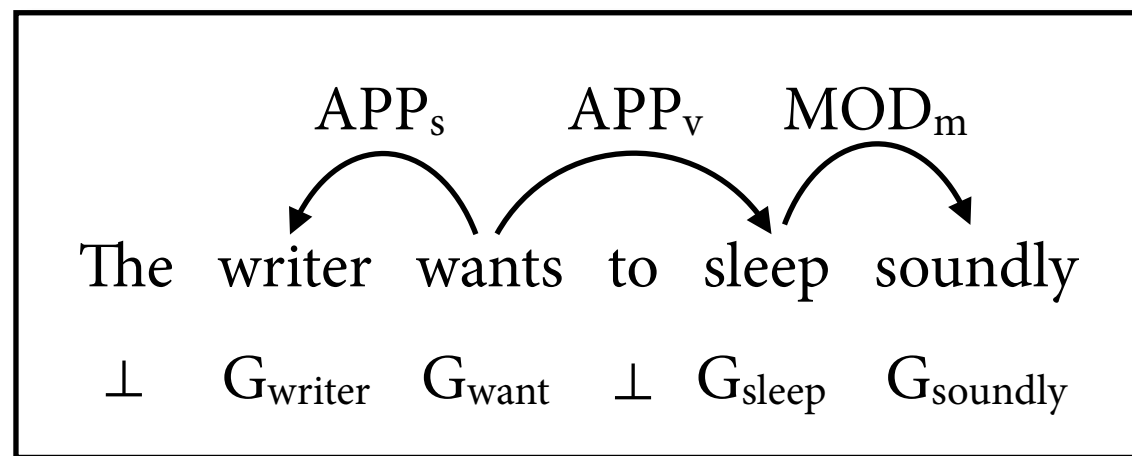
Relation Identification: add all edges with positive weight; then repeatedly add least negative edge that connects subgraphs.

Issues with JAMR

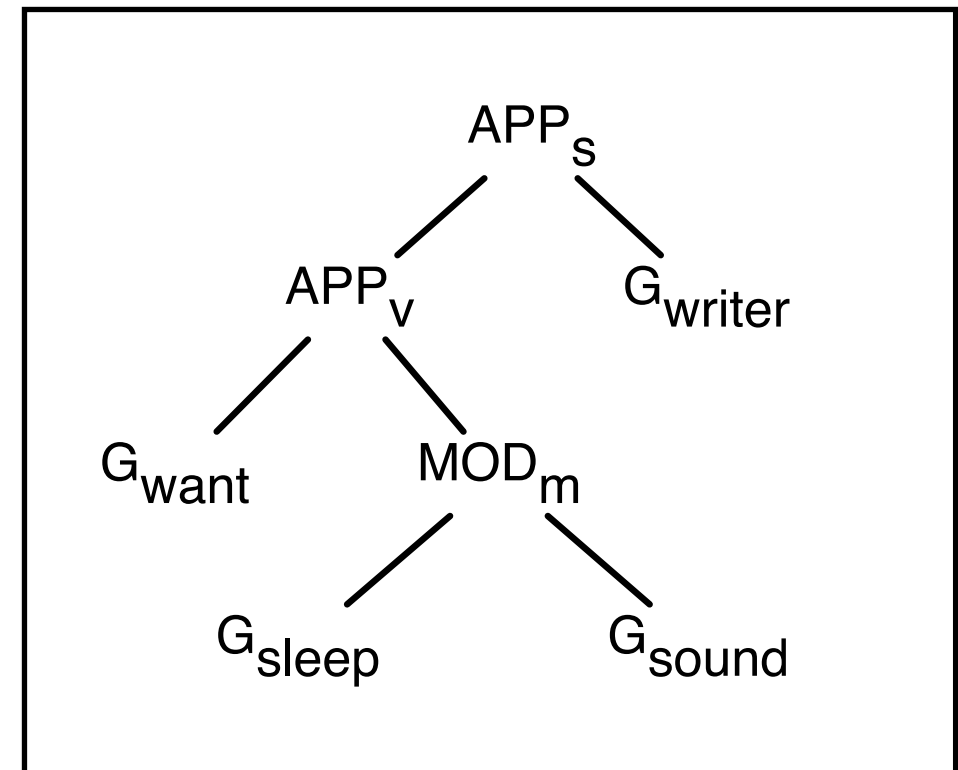
- JAMR can draw edge between any two nodes; syntactic structure of sentence used only indirectly.
- Semantic representations for words don't know anything about their semantic arguments.
- Edges for control verbs added arbitrarily, not because linguistic principle of control discovered.
- No notion of compositionality!

Compositional semantic parsing


Learn to predict a (semantic) dependency tree,
which then decodes into an AM term, which evaluates to a graph.



AM dependency tree





Term over AM graph algebra



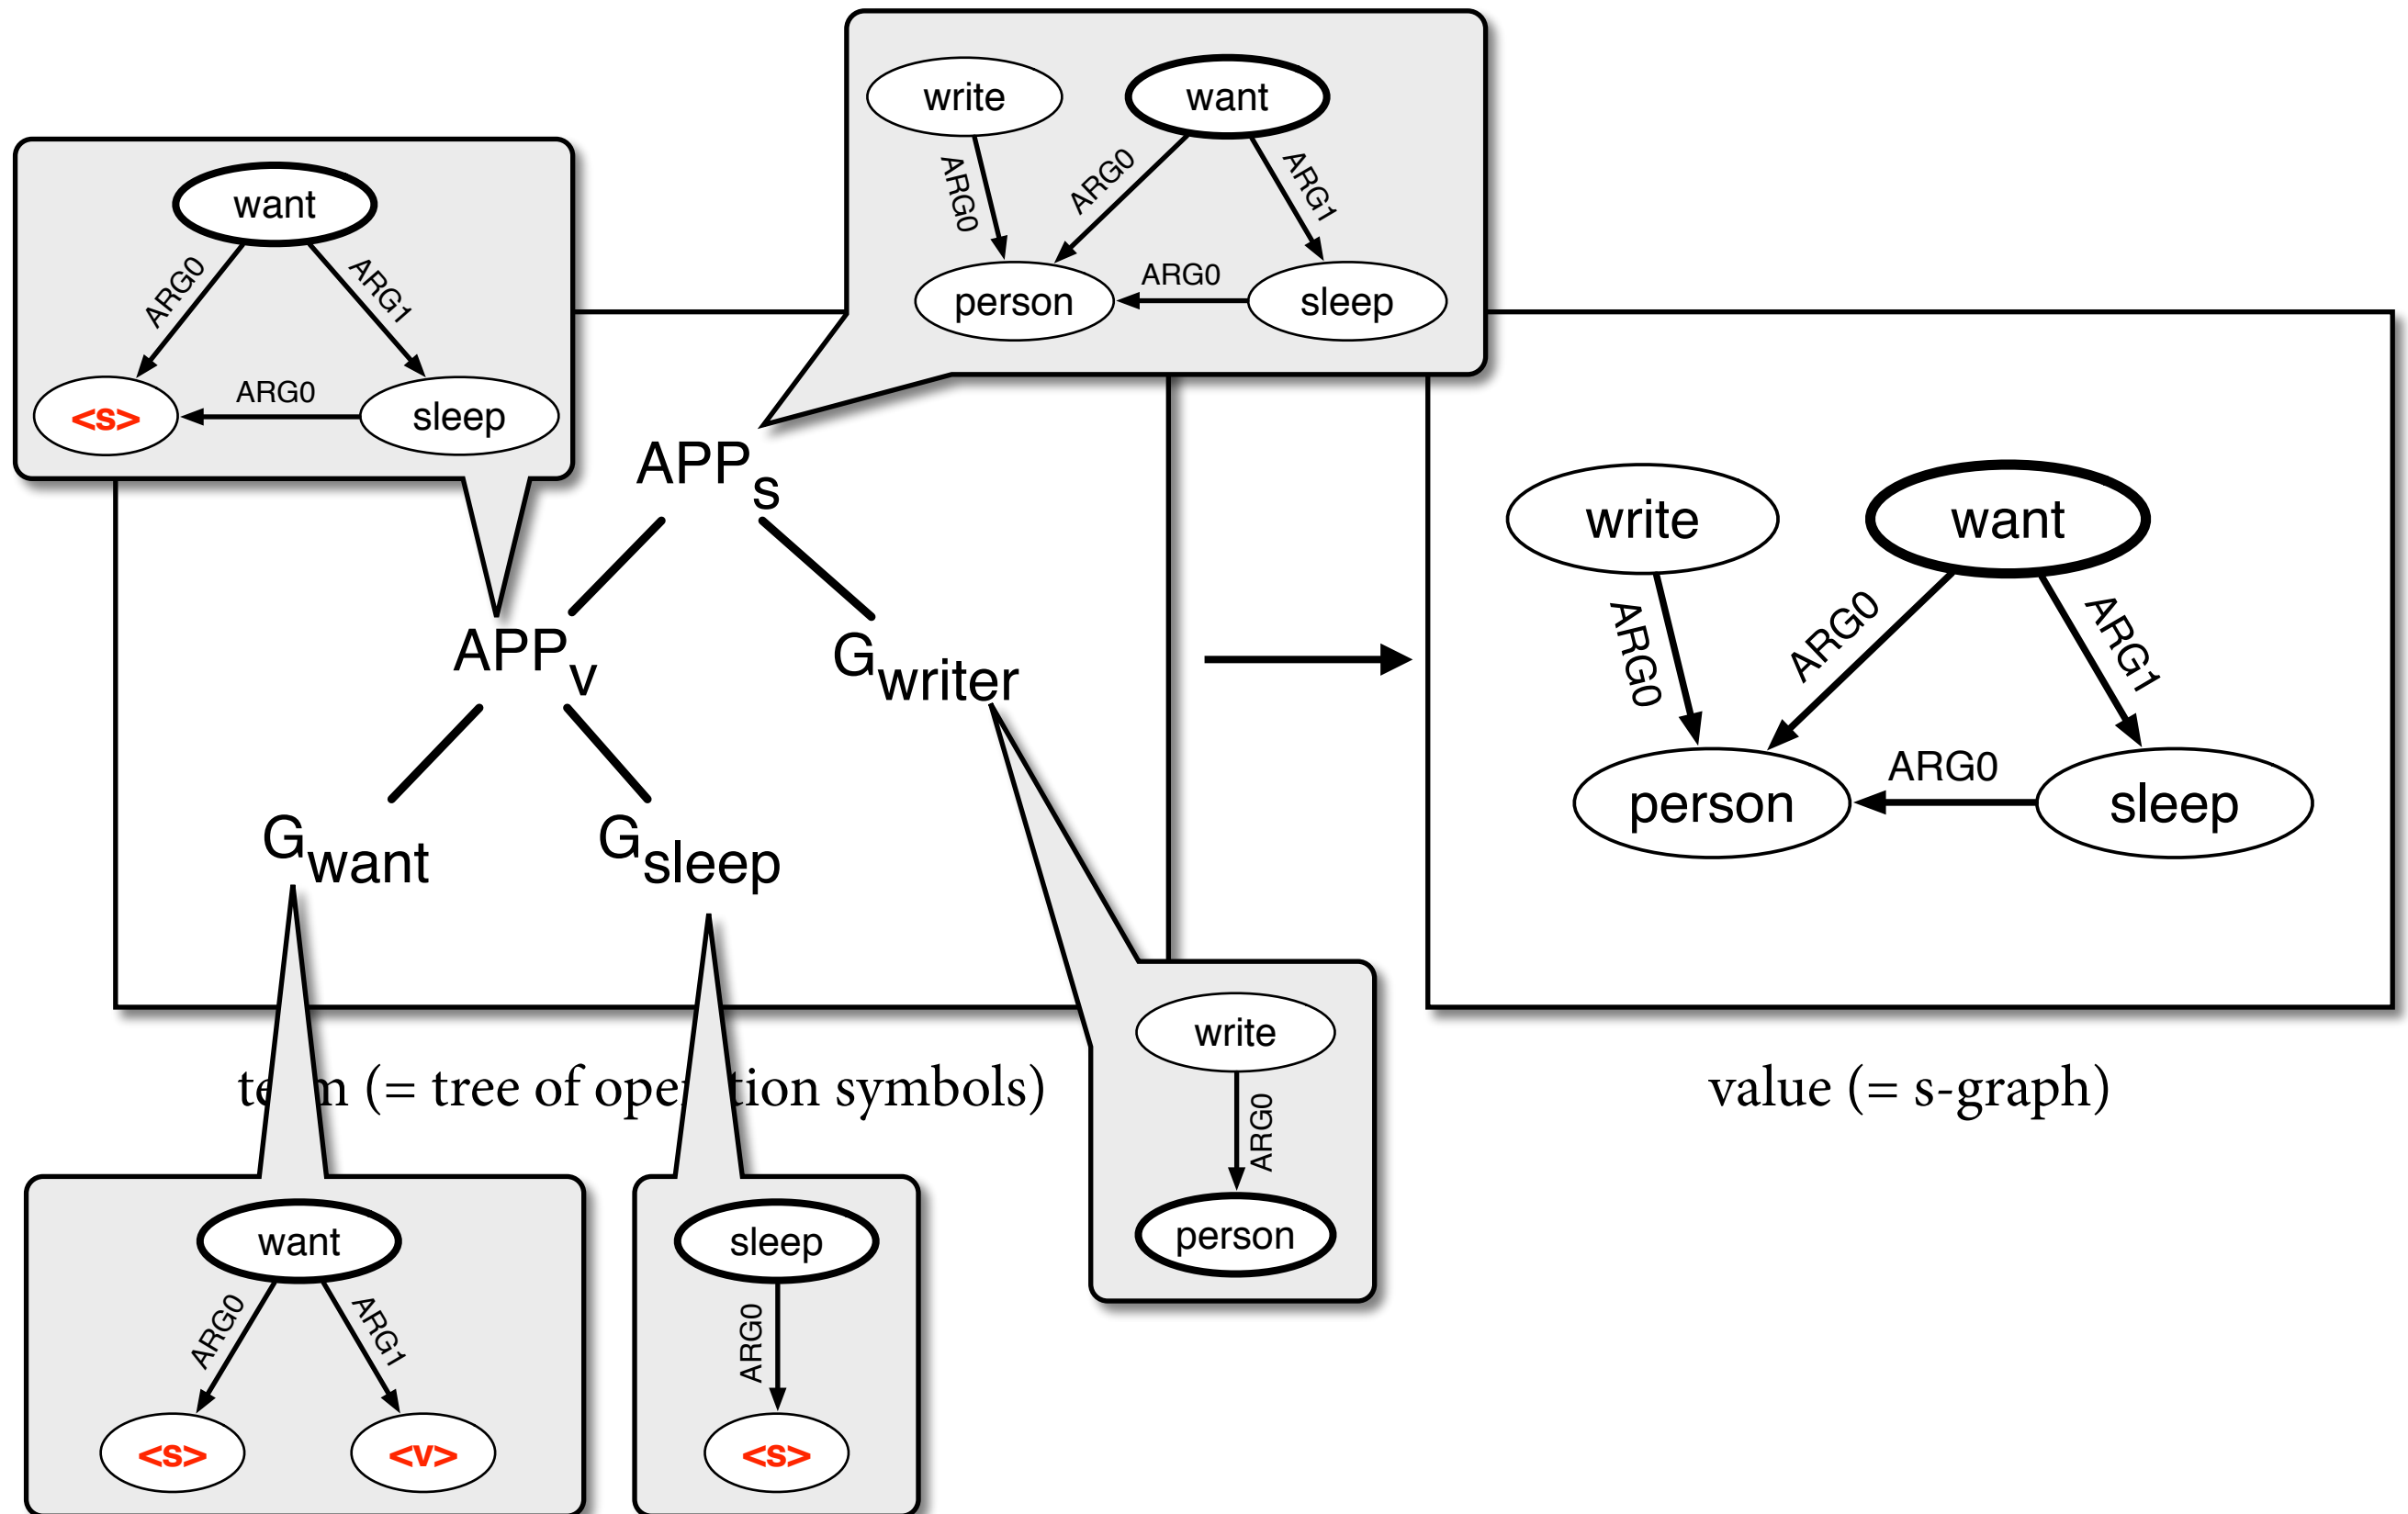
graph (e.g. AMR)

AM algebra

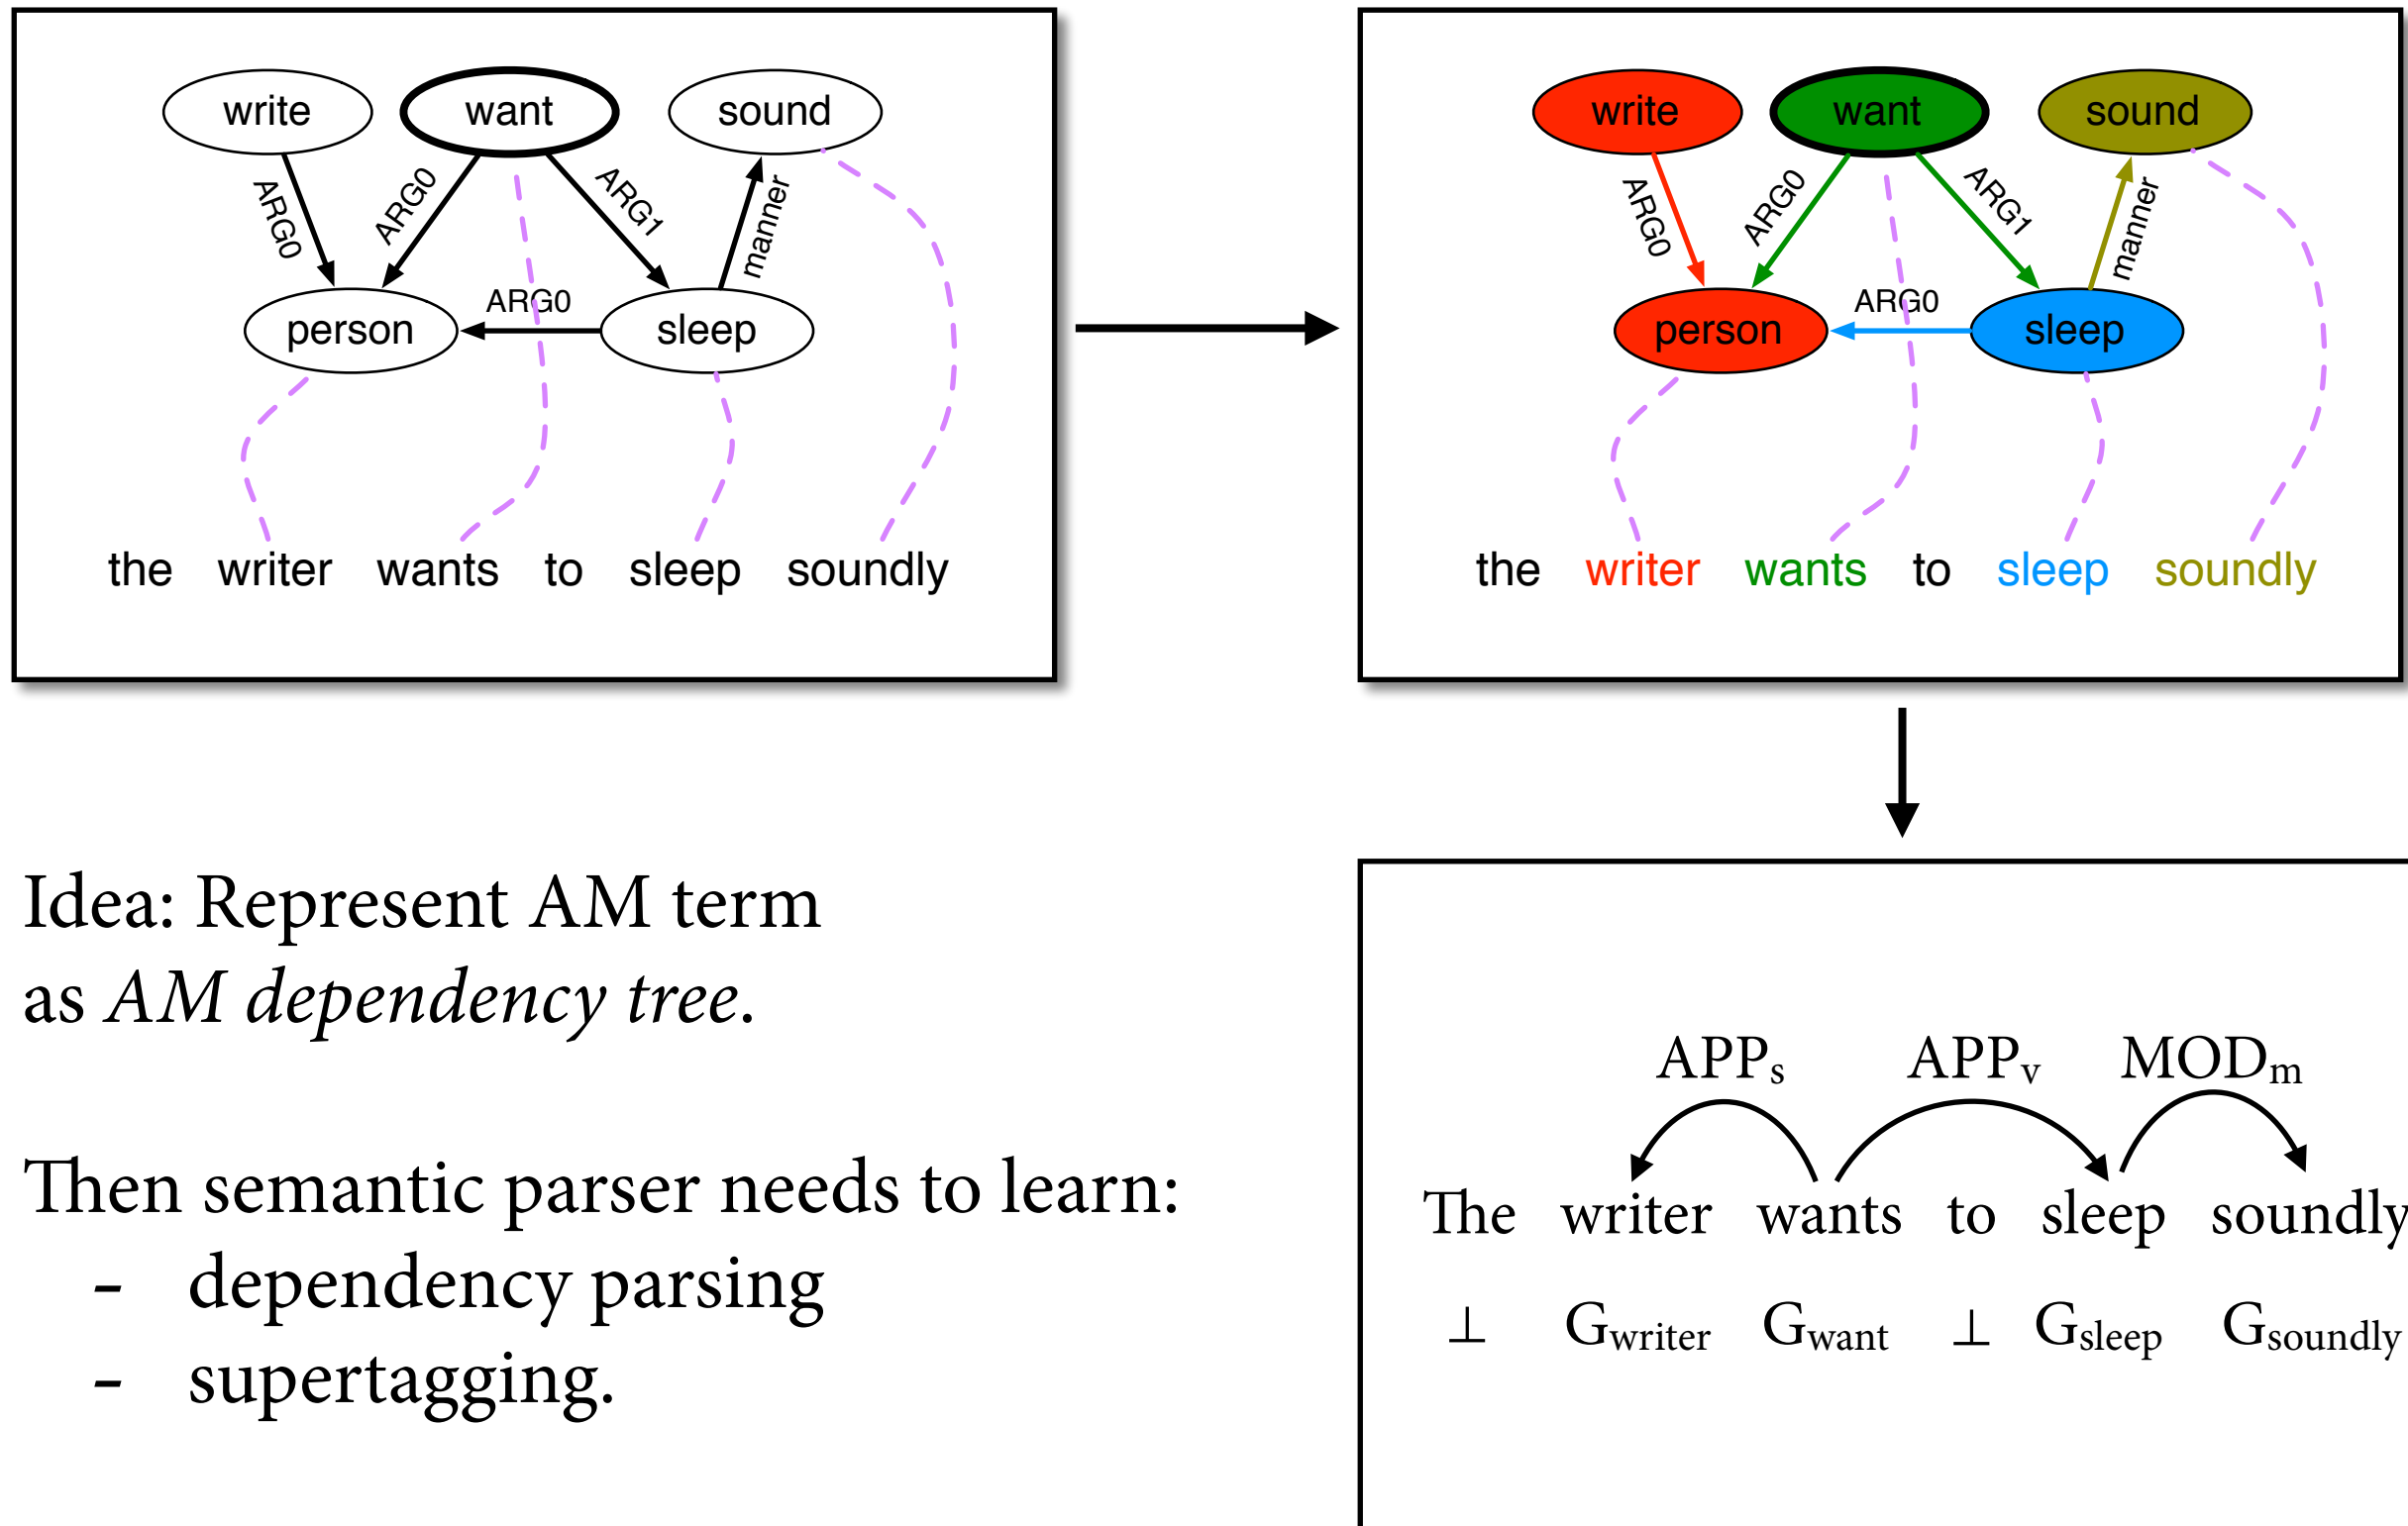
Two operations for combining s-graphs:
Apply (= head + complement), Modify (= head + modifier).



AM terms



Converting training data



Idea: Represent AM term
as *AM dependency tree*.

Then semantic parser needs to learn:

- dependency parsing
- supertagging.

Parsing

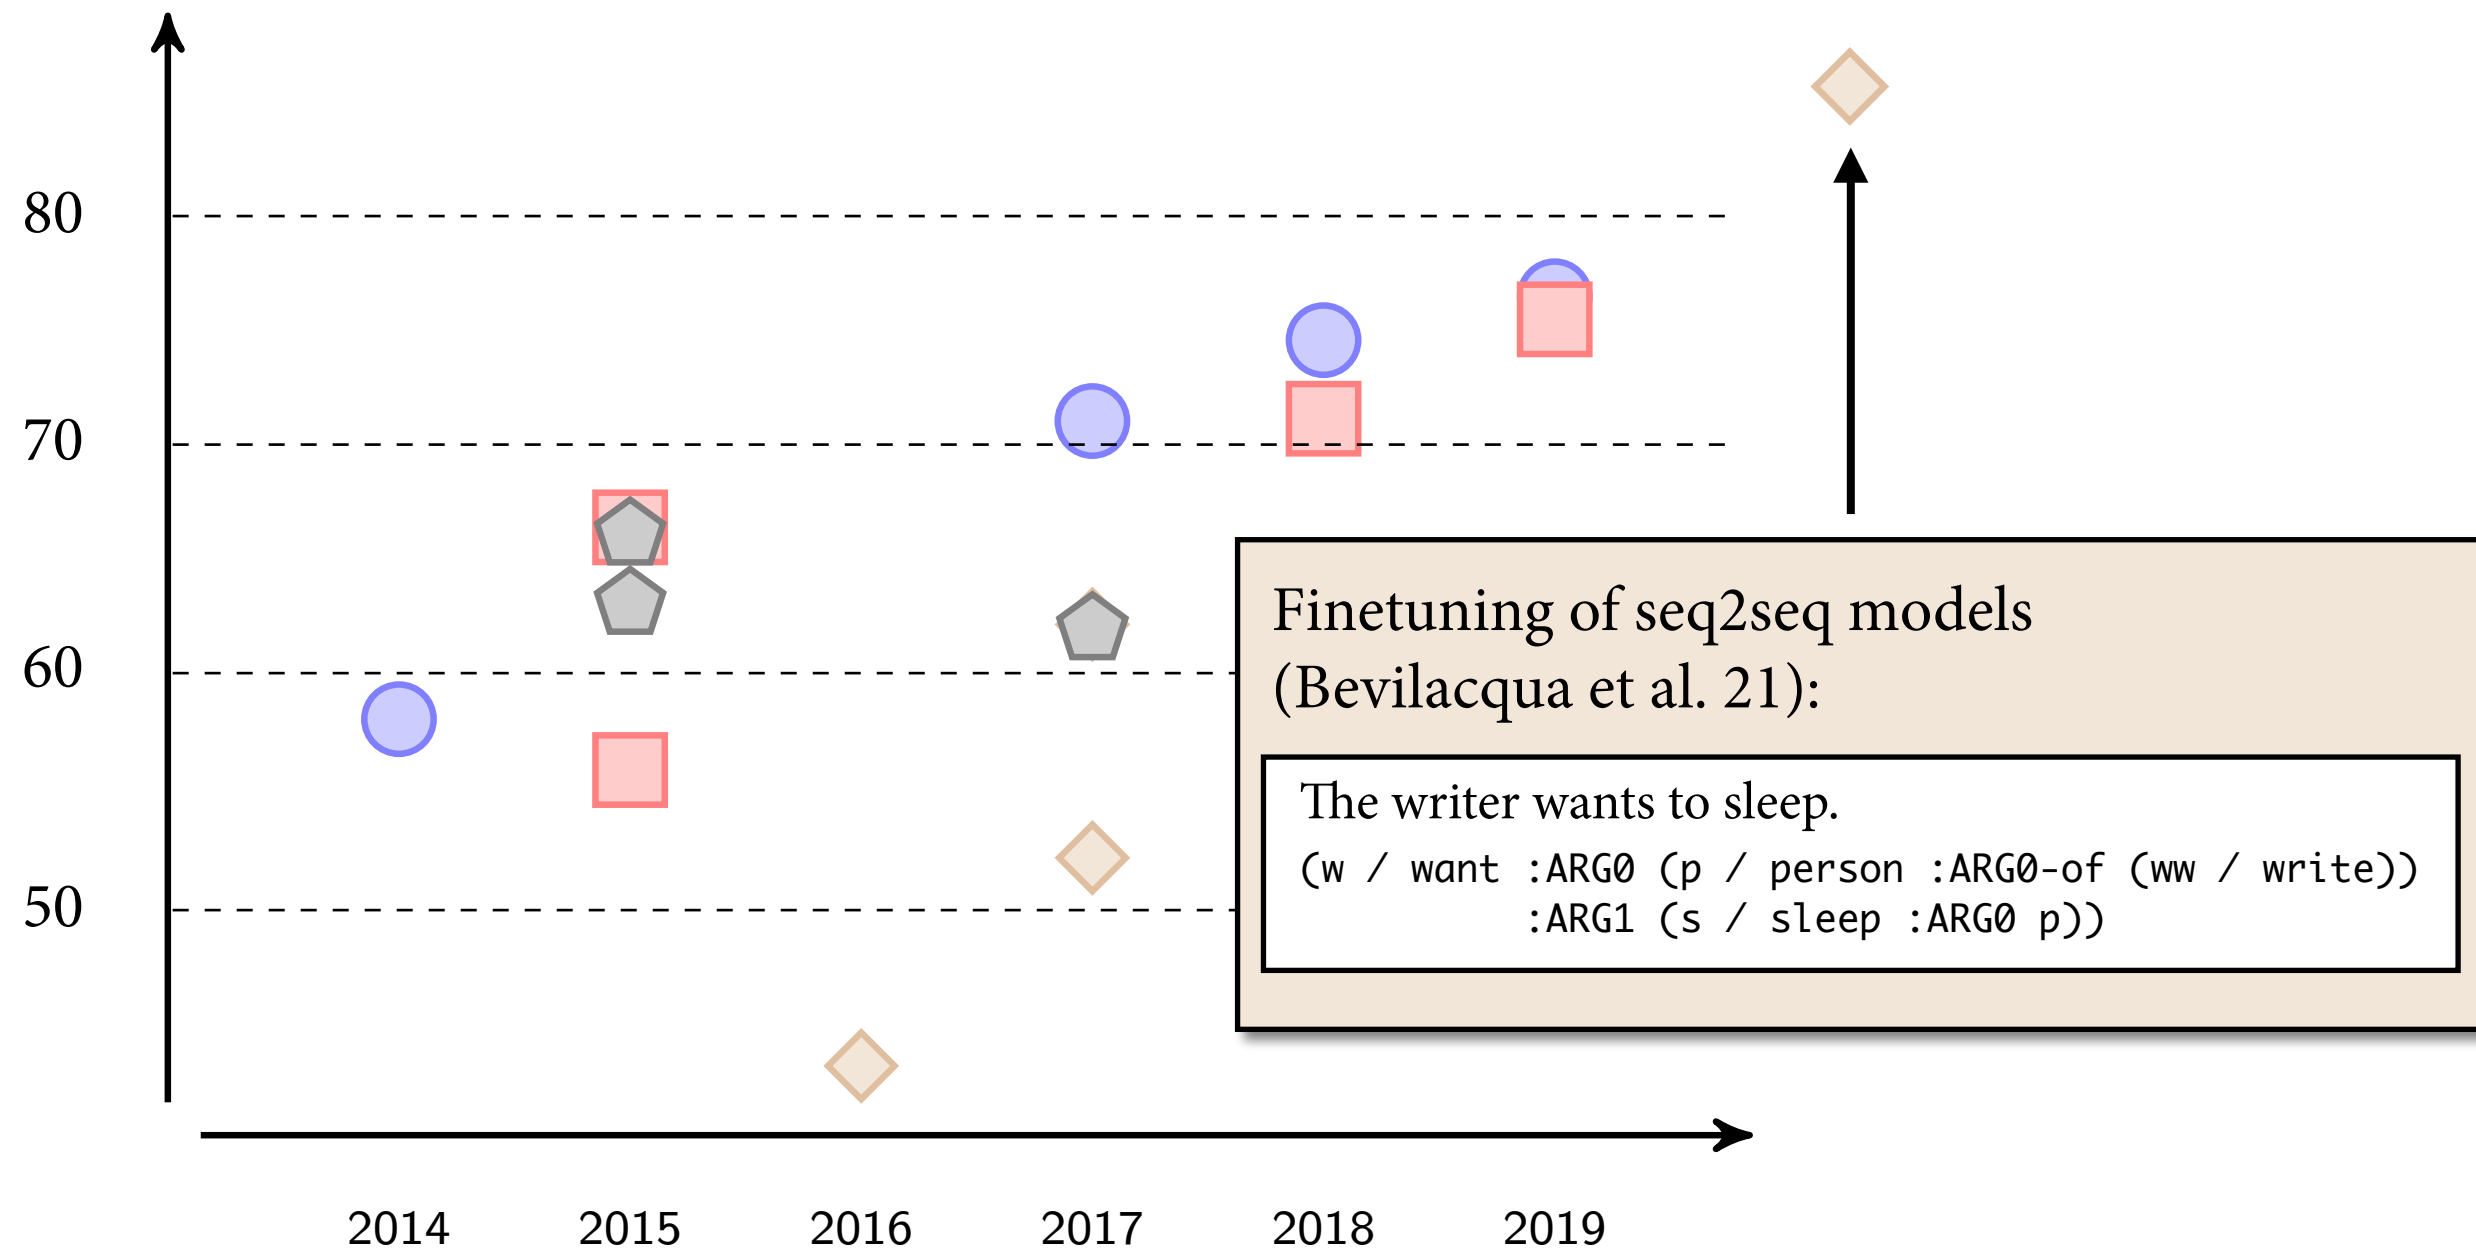
- Convert (string, graph) training data into (string, supertags + dependencies) training data.
- Train neural supertagger + dependency parser to assign scores to supertags + dependencies.
 - ▶ easier than predicting the whole graph; compositional!
- At evaluation time, compute highest-scoring well-typed dependency tree.
 - ▶ well-typedness requirement makes this NP-complete
 - ▶ solve approximately with CKY-style parsing algorithm

Parsing across graphbanks


	DM		PAS		PSD		EDS		AMR 2015	AMR 2017
	id F	ood F	id F	ood F	id F	ood F	Smatch F	EDM	Smatch F	Smatch F
Groschwitz et al. (2018)	-	-	-	-	-	-	-	-	70.2	71.0
Lyu and Titov (2018)	-	-	-	-	-	-	-	-	73.7	74.4 ± 0.16
Zhang et al. (2019)	-	-	-	-	-	-	-	-	-	76.3 ± 0.1
Peng et al. (2017) Basic	89.4	84.5	92.2	88.3	77.6	75.3	-	-	-	-
Dozat and Manning (2018)	93.7	88.9	94.0	90.8	81.0	79.4	-	-	-	-
Buys and Blunsom (2017)	-	-	-	-	-	-	85.5	85.9	60.1	-
Chen et al. (2018)	-	-	-	-	-	-	90.9 ^{1,2}	90.4 ¹	-	-
This paper (GloVe)	90.4 ± 0.2	84.3 ± 0.2	91.4 ± 0.1	86.6 ± 0.1	78.1 ± 0.2	74.5 ± 0.2	87.6 ± 0.1	82.5 ± 0.1	69.2 ± 0.4	70.7 ± 0.2
This paper (BERT)	93.9 ± 0.1	90.3 ± 0.1	94.5 ± 0.1	92.5 ± 0.1	82.0 ± 0.1	81.5 ± 0.3	90.1 ± 0.1	84.9 ± 0.1	74.3 ± 0.2	75.3 ± 0.2
Peng et al. (2017) Freda1	90.0	84.9	92.3	88.3	78.1	75.8	-	-	-	-
Peng et al. (2017) Freda3	90.4	85.3	92.7	89.0	78.5	76.4	-	-	-	-
This paper, MTL (GloVe)	91.2 ± 0.1	85.7 ± 0.0	92.2 ± 0.2	88.0 ± 0.3	78.9 ± 0.3	76.2 ± 0.4	88.2 ± 0.1	83.3 ± 0.1	(70.4) ³ ± 0.2	71.2 ± 0.2
This paper, MTL (BERT)	94.1 ± 0.1	90.5 ± 0.1	94.7 ± 0.1	92.8 ± 0.1	82.1 ± 0.2	81.6 ± 0.1	90.4 ± 0.1	85.2 ± 0.1	(74.5) ³ ± 0.1	75.3 ± 0.1

- First semantic parser that does well across all six major graphbanks.
- Established new states of the art through use of pretrained BERT embeddings.
- Small improvements through multi-task learning on multiple graphbanks.
- Transition-based parser can be extremely fast ($\sim 10k$ tokens/sec) and accurate.

Progress in semantic parsing





Seq2seq Semantic Parsing



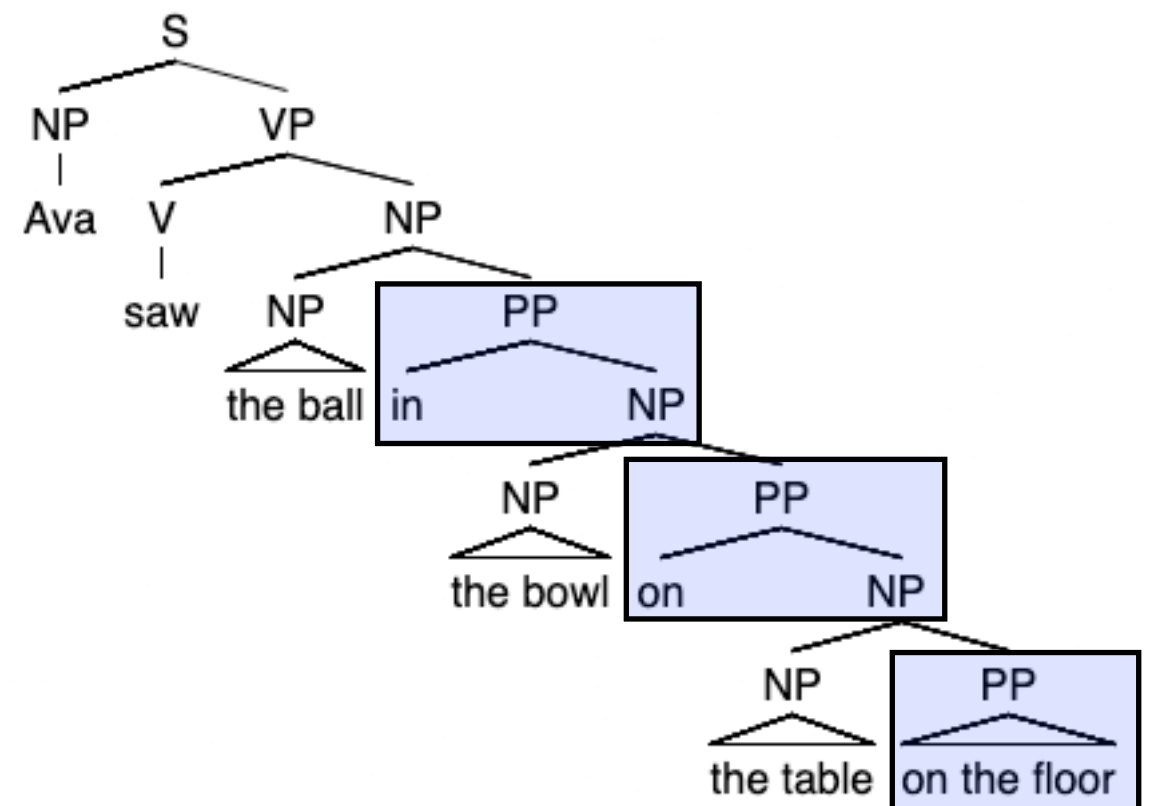
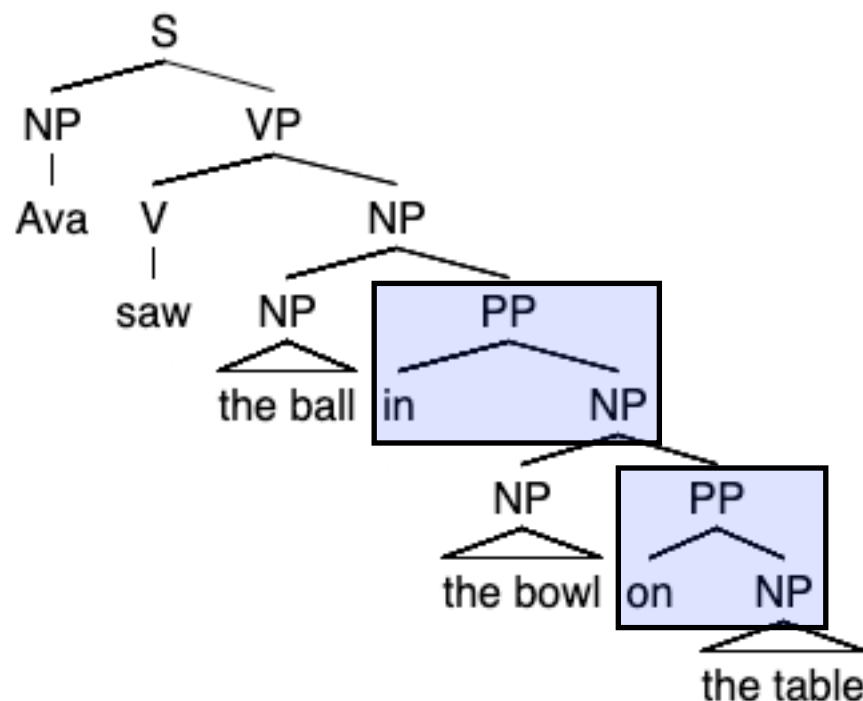
Compositional generalization

Structural: observe PPs in objects, generalize to PPs in subjects ("objPP to subjPP")



Noah ate the cake on the table.
*cake(x_3); *plate(x_6);
eat.agent(x_1 , Noah) \wedge eat.theme(x_1 , x_3)
 \wedge cake.nmod.on(x_3 , x_6)

Structural: observe PPs up to recursion depth 2, generalize to recursion depth 3-12



(COGS corpus, Kim & Linzen 20)

Compositional helps on COGS

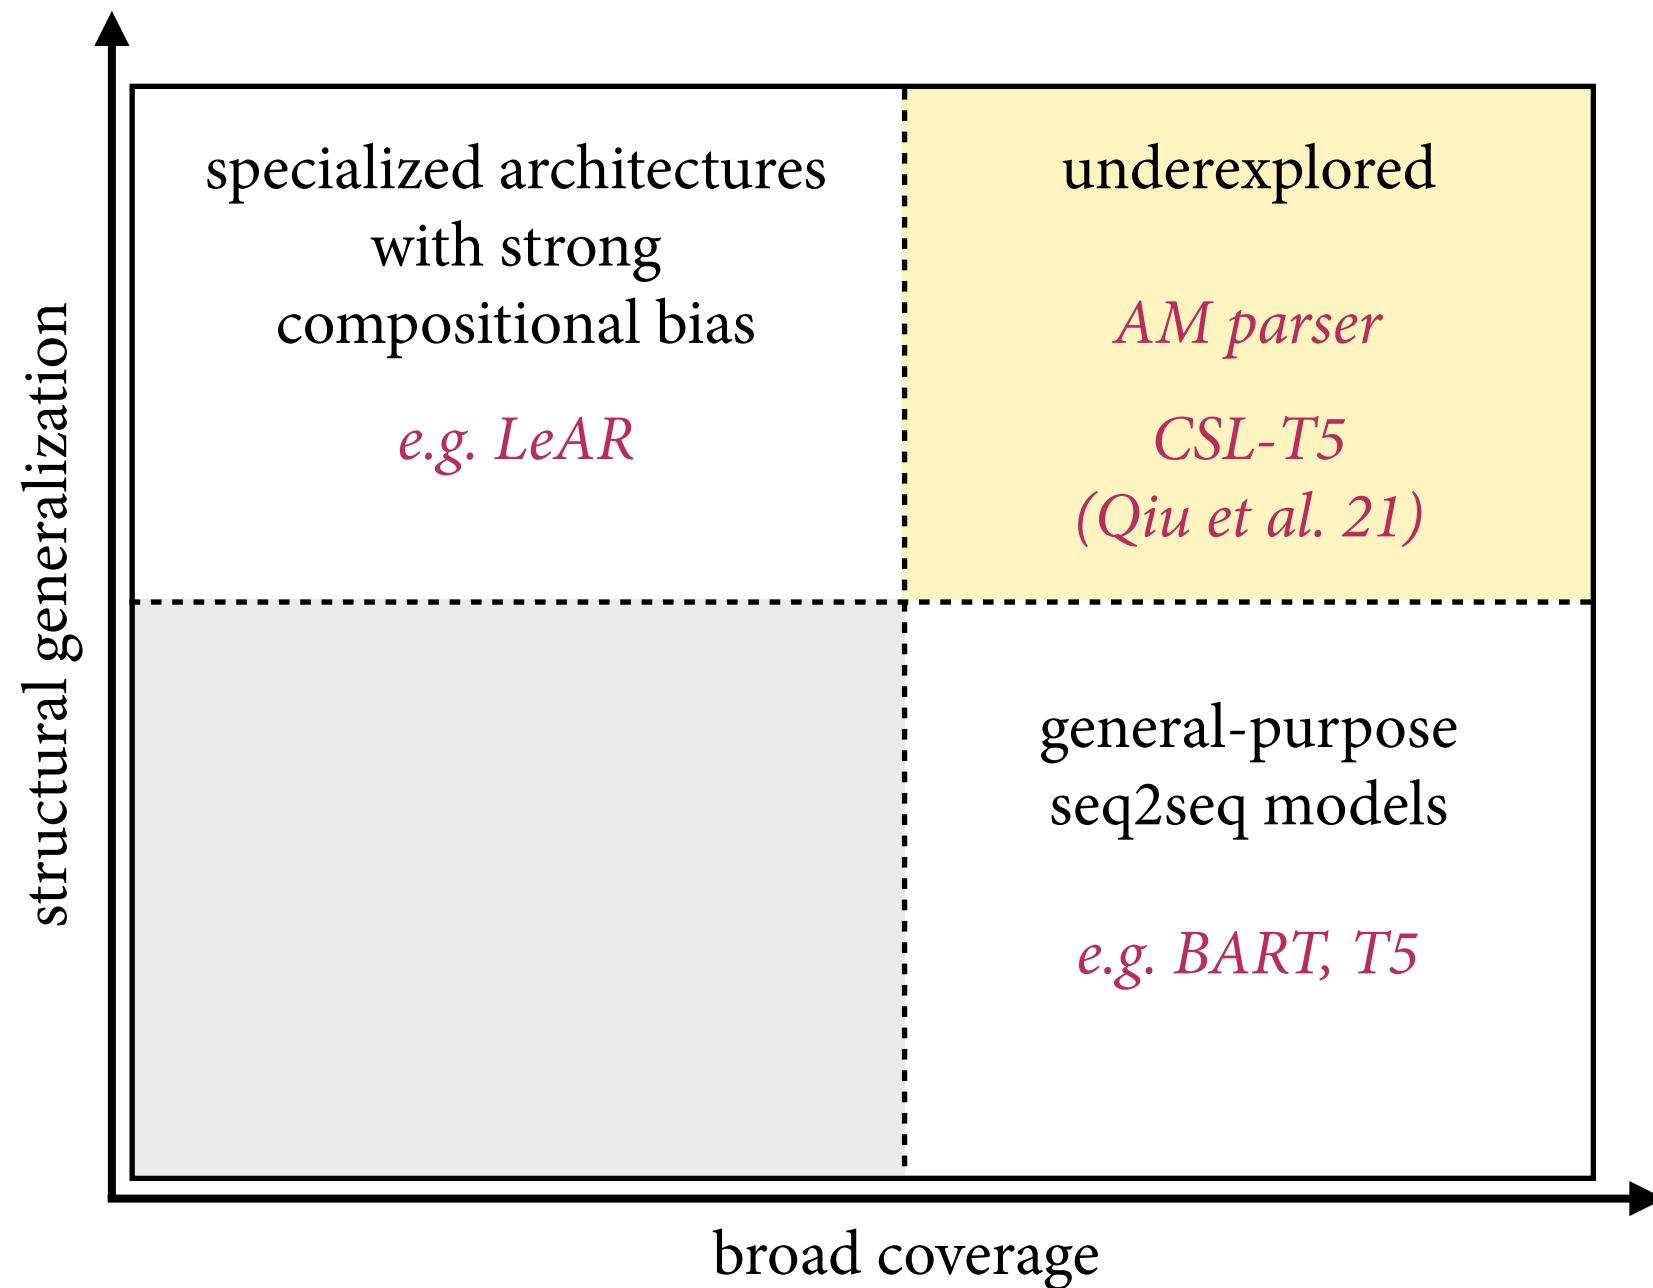
Model Class	Model	STRUCT			LEX	Overall
		Obj to Subj PP	CP recursion	PP recursion	all 18 other types	
seq2seq	BART	0	0	12	91	79
	BART+syn	0	5	8	93	80
	T5	0	0	9	97	83
	Kim and Linzen 2020	0	0	0	73	63
	Akyürek and Andreas 2021	0	0	1	96	82
	Zheng and Lapata 2022	0	12	39	99	89
	Conklin et al. 2021	0	0	0	88	75
	Csordás et al. 2021	0	0	0	95	81
	Qiu et al. 2021 *	100	100	100	100	100
structure-aware	Liu et al. 2021	93	100	99	99	99
	Weißenhorn et al. 2022	78	100	99	100	98

AM parser →

*) Can still see this as a neurosymbolic model (uses compositional data augmentation); achieves perfect accuracy, up to depth seen in the augmented data.

(Weißenhorn, Donatelli, K., *SEM 2022; Yao & K., EMNLP 2022)

Semantic parsing tasks



(picture adapted from Shaw et al., ACL 2021)

Conclusion

- Challenge in compositional semantic construction:
Where do we get large-scale grammars?
- Semantic parsing: Learn such grammars from corpora with semantic annotations.
 - ▶ GeoQuery: small corpus of trees
 - ▶ graphbanks, e.g. AMRBank
 - ▶ also parsing into programs and other formal languages